



**Please read pinned Ed posts!**

# **From Bits through Integers**

15-213/14-513/15-513:  
Introduction to Computer Systems

2<sup>nd</sup> Lecture, Spring 2026

# Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
- Byte Ordering

CSAPP 2.1

CSAPP 2.2

CSAPP 2.3

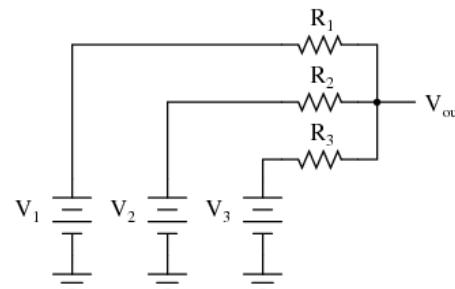
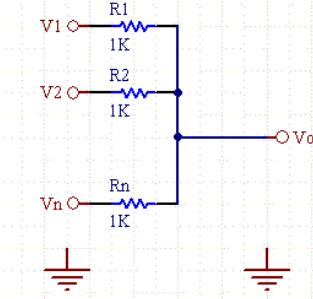
CSAPP 2.1.3

# Analog Computers

■ Before digital computers there were analog computers.

■ Consider a couple of simple analog computers:

- A simple circuit can allow one to adjust voltages using variable resistors and measure the output using a volt meter:
- A simple network of adjustable parallel resistors can allow one to find the average of the inputs.

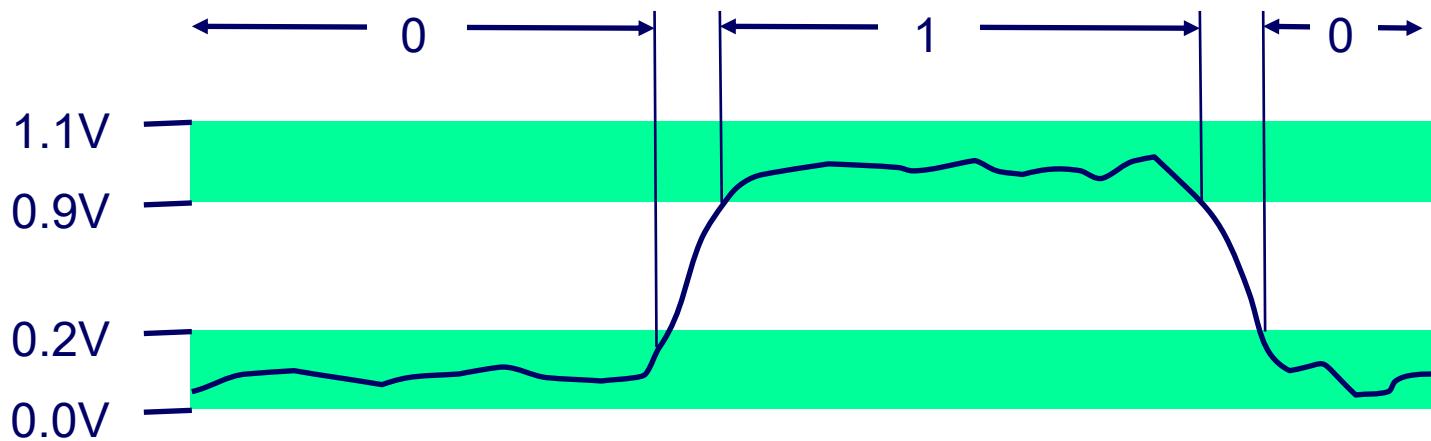


<https://www.daycounter.com/Calculators/Voltage-Summer/Voltage-Summer-Calculator.php>

<https://www.quora.com/What-is-the-most-basic-voltage-adder-circuit-without-a-transistor-op-amp-and-any-external-supply>

# Needing Less Accuracy+Precision is Easier

- We don't try to measure exactly
  - We just ask, is it high enough to be "On", or
  - Is it low enough to be "Off".
- We have two states, so we have a binary, or 2-ary, system.
  - We represent these states as 0 and 1
- Now we can easily interpret, communicate, and duplicate signals well enough to know what they mean.



# Binary Representation

- **Binary representation -> binary, base-2, numbering system**
  - 0 represents 0
  - 1 represents 1
  - Each “place” represents a power of two, exactly as each place in our usual “base 10”, 10-ary numbering system represents a power of ten
- **Encoding/interpreting sets of bits can represent...**
  - Operations to be executed by the processor
  - numbers
  - enumerable things, such as text characters
  - more!
- **If we can assign it to a discrete number, we can represent it in binary**

# Binary Representation: Simple Numbers

- Example: we can count in binary, a base-2 numbering system

- 000, 001, 010, 011, 100, 101, 110, 111, ...
  - $000 = 0*2^2 + 0*2^1 + 0*2^0 = 0$  (in decimal)
  - $001 = 0*2^2 + 0*2^1 + 1*2^0 = 1$  (in decimal)
  - $010 = 0*2^2 + 1*2^1 + 0*2^0 = 2$  (in decimal)
  - $011 = 0*2^2 + 1*2^1 + 1*2^0 = 3$  (in decimal)
  - Etc.

- For reference, consider some base-10 examples:

- $000 = 0*10^2 + 0*10^1 + 0*10^0$
- $001 = 0*10^2 + 0*10^1 + 1*10^0$
- $357 = 3*10^2 + 5*10^1 + 7*10^0$

# Hexadecimal and Octal

- Writing numbers in binary takes too many digits
- Want a way to represent numbers more densely such that fewer digits are required
  - But also so it is easy to get at the bits that we want
- Any power-of-two base provides this property
  - Octal (base 8) and hexadecimal (base 16) are closest to our familiar base-10.
  - Each has been used by “computer people” over time
  - Hexadecimal is often preferred because it is denser.

# Hexadecimal

## ■ Hexadecimal $00_{16}$ to $FF_{16}$

- Base 16 number representation
- Use characters ‘0’ to ‘9’ and ‘A’ to ‘F’

## ■ Consider $1A2B$ in Hexadecimal:

- $1*16^3 + A*16^2 + 2*16^1 + B*16^0$
- $1*16^3 + 10*16^2 + 2*16^1 + 11*16^0 = 6699$  (decimal)

- The C Language prefixes hexadecimal numbers with “0x” so they aren’t confused with decimal numbers
- Write  $FA1D37B_{16}$  in C as

- `0xFA1D37B`
- `0xfa1d37b`

$15213 : \underline{0011} \quad \underline{1011} \quad \underline{0110} \quad \underline{1101}$

3            B            6            D

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
- Byte Ordering

# Boolean Algebra

## ■ Developed by George Boole in 19th Century

- Algebraic representation of logic
  - Encode “True” as 1 and “False” as 0

## And

- $A \& B = 1$  when both  $A=1$  and  $B=1$

&	0	1
0	0	0
1	0	1

## Or

- $A | B = 1$  when either  $A=1$  or  $B=1$

	0	1
0	0	1
1	1	1

## Not

- $\sim A = 1$  when  $A=0$

$\sim$	
0	1
1	0

## Exclusive-Or (Xor)

- $A \wedge B = 1$  when either  $A=1$  or  $B=1$ , but not both

$\wedge$	0	1
0	0	1
1	1	0

# The Boolean Operators Work On ALL bits

- Operations applied bitwise

$$\begin{array}{rcl} \begin{array}{c} 01101001 \\ \& 01010101 \\ \hline 01000001 \end{array} & \begin{array}{c} 01101001 \\ | 01010101 \\ \hline 01111101 \end{array} & \begin{array}{c} 01101001 \\ ^ 01010101 \\ \hline 00111100 \end{array} \\ & & \begin{array}{c} \sim 01010101 \\ \hline 10101010 \end{array} \end{array}$$

- All of the Properties of Boolean Algebra Apply

# Example: Representing & Manipulating Sets

## ■ Representation

- Width w bit vector represents subsets of  $\{0, \dots, w-1\}$
- $a_j = 1$  if  $j \in A$

- 01101001             $\{0, 3, 5, 6\}$

- ~~76543210~~

- 01010101             $\{0, 2, 4, 6\}$

- ~~76543210~~

## ■ Operations

■ & Intersection	01000001	$\{0, 6\}$
■   Union	01111101	$\{0, 2, 3, 4, 5, 6\}$
■ ^ Symmetric difference	00111100	$\{2, 3, 4, 5\}$
■ ~ Complement	10101010	$\{1, 3, 5, 7\}$

# Bit-Level Operations in C

## Operations &, |, ~, ^ Available in C

- Apply to any “integral” data type
  - long, int, short, char, unsigned
- View arguments as bit vectors
- Arguments applied bit-wise

## Examples (Char data type)

- $\sim 0x41 \rightarrow 0xBE$ 
  -
- $\sim 0x00 \rightarrow 0xFF$ 
  -
- $0x69 \& 0x55 \rightarrow 0x41$ 
  -
- $0x69 | 0x55 \rightarrow 0x7D$ 
  -

	Hex	Decimal	Binary
0	0	0000	
1	1	0001	
2	2	0010	
3	3	0011	
4	4	0100	
5	5	0101	
6	6	0110	
7	7	0111	
8	8	1000	
9	9	1001	
A	10	1010	
B	11	1011	
C	12	1100	
D	13	1101	
E	14	1110	
F	15	1111	

# Bit-Level Operations in C

## Operations &, |, ~, ^ Available in C

- Apply to any “integral” data type
  - long, int, short, char, unsigned
- View arguments as bit vectors
- Arguments applied bit-wise

## Examples (Char data type)

- $\sim 0x41 \rightarrow 0xBE$ 
  - $\sim 01000001_2 \rightarrow 10111110_2$
- $\sim 0x00 \rightarrow 0xFF$ 
  - $\sim 00000000_2 \rightarrow 11111111_2$
- $0x69 \& 0x55 \rightarrow 0x41$ 
  - $01101001_2 \& 01010101_2 \rightarrow 01000001_2$
- $0x69 | 0x55 \rightarrow 0x7D$ 
  - $01101001_2 | 01010101_2 \rightarrow 01111101_2$

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

# Bit-Level Operations in C

## ■ Operations &, |, ~, ^ Available in C

- Apply to any “integral” data type
  - long, int, short, char, unsigned
- View arguments as bit vectors
- Arguments applied bit-wise

## ■ Examples (Char data type)

- $\sim 0x41 \rightarrow 1011\ 1110$

- $\sim 0x00 \rightarrow 1111\ 1111$

- $0x69 \ \& \ 0x55:$

0110 1001

& 0101 0101

$0x69 \mid 0x55:$

0110 1001

| 0101 0101

---

0100 0001

---

0111 1101

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

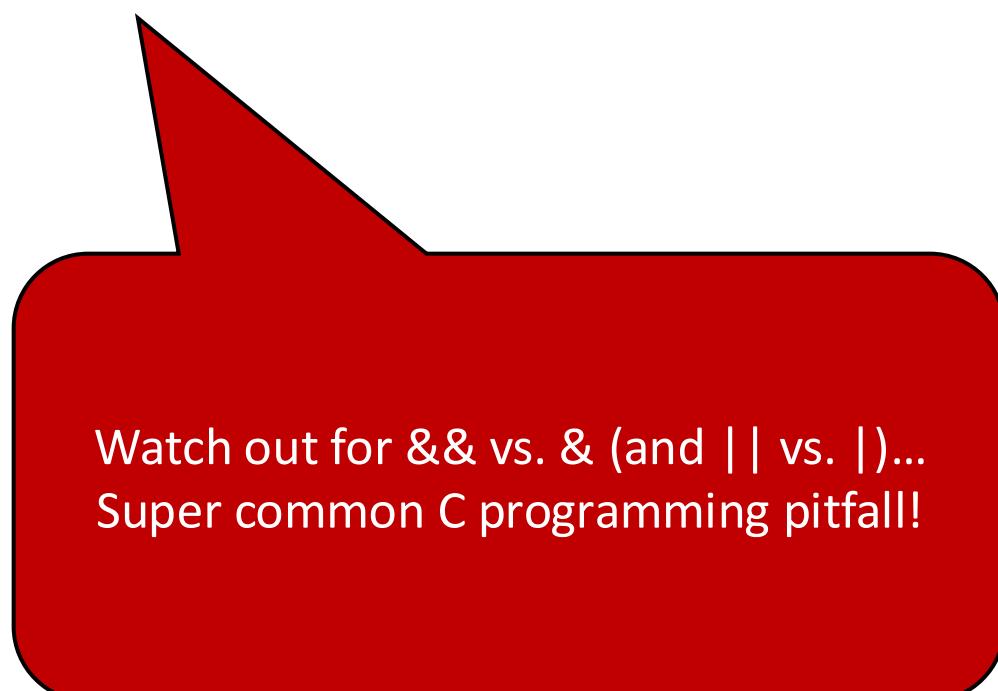
# Contrast: Logic Operations in C

## ■ Contrast to Bit-Level Operators

- Logic Operations: `&&`, `||`, `!`
  - View 0 as “False”
  - Anything nonzero as “True”
  - Always return 0 or 1
  - Early termination

## ■ Examples (char data type)

- `!0x41 → 0x00`
- `!0x00 → 0x01`
- `!!0x41 → 0x01`
- `0x69 && 0x55 → 0x01`
- `0x69 || 0x55 → 0x01`
- `p && *p` (avoids null pointer access)



Watch out for `&&` vs. `&` (and `||` vs. `|`)...  
Super common C programming pitfall!

# Shift Operations

## ■ Left Shift: $x \ll y$

- Shift bit-vector **x** left **y** positions
  - Throw away extra bits on left
  - Fill with 0's on right

## ■ Right Shift: $x \gg y$

- Shift bit-vector **x** right **y** positions
  - Throw away extra bits on right
- Logical shift
  - Fill with 0's on left
- Arithmetic shift
  - Replicate most significant bit on left

Argument <b>x</b>	
$\ll 3$	
Log. $\gg 2$	
Arith. $\gg 2$	

Argument <b>x</b>	
$\ll 3$	
Log. $\gg 2$	
Arith. $\gg 2$	

## ■ Undefined Behavior

- Shift amount  $< 0$  or  $\geq$  word size

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting

# Binary Number Lines

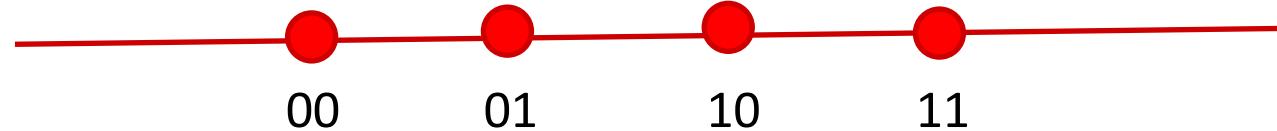
- In binary, the number of bits in the data type size determines the number of points on the number line.
  - We can assign the points any meaning we'd like

- Consider the following examples:

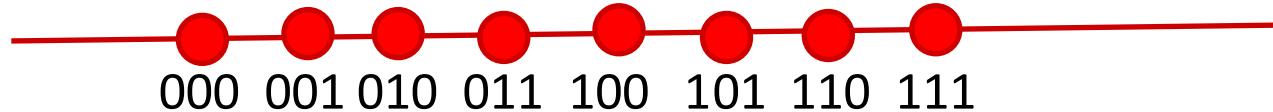
- 1 bit number line



- 2 bit number line

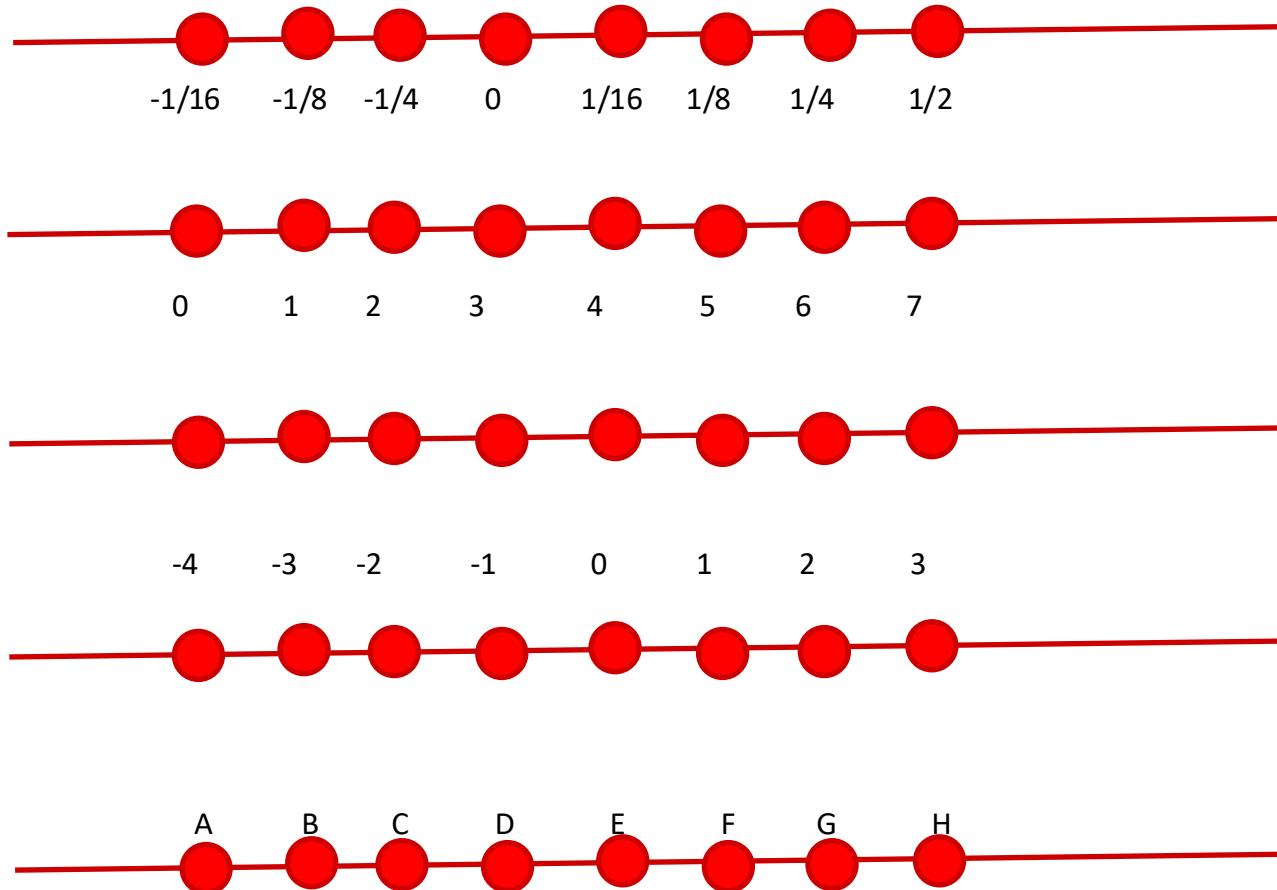


- 3 bit number line



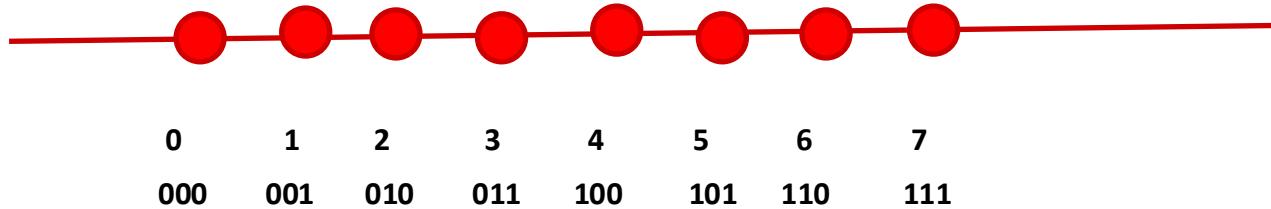
# Some Purely Imaginary Examples

## ■ 3 bit number line



# Overflow

- Let's consider a simple 3 digit number line:



- What happens if we add 1 to 7?
  - In other words, what happens if we add 1 to 111?
- $111 + 001 = 1\ 000$ 
  - But, we only get 3 bits – so we lose the leading 1.
  - This is called overflow
- The result is 000

# Modulus Arithmetic

- Let's explore this idea of overflow some more

- $111 + 001 = 1\ 000 = 000$
- $111 + 010 = 1\ 001 = 001$
- $111 + 011 = 1\ 010 = 010$
- $111 + 100 = 1\ 011 = 011$
- ...
- $111 + 110 = 1\ 101 = 101$
- $111 + 111 = 1\ 110 = 110$

- So, arithmetic “wraps around” when it gets “too positive”

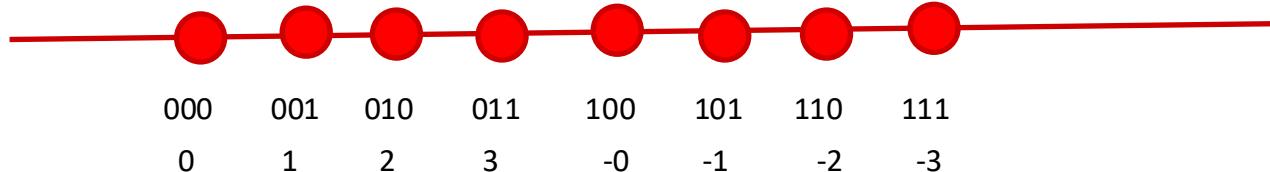
# Unsigned and Non-Negative Integers

- We'll use the term “ints” to mean the finite set of integer numbers that we can represent on a number line enumerated by some fixed number of bits, i.e. *bit width*.
- We normally represent unsigned and non-negative int using simple binary as we have already discussed
  - An “unsigned” int is any int on a number line, e.g. of a data type, that doesn't contain any negative numbers
  - A non-negative number is a number greater than or equal to ( $\geq$ ) 0 on a number line, e.g. of a data type, that does contain negative numbers

# How to represent negative Numbers?

- We could use the leading bit as a *sign bit*.

- 0 means non-negative
- 1 means negative



- Benefits

- represents negative and non-negative numbers
- 0 represents 0

- Drawbacks

- There is a -0, which is the same as 0, except that it is different
- How to add such numbers  $1 + -1$  should equal 0
  - But, by simple math,  $001 + 101 = 110$ , which is -2?

# A Trick!

## ■ Let's start with three ideas:

- 1 should be represented as 1
- $-1 + 1 = 0$
- We want addition to work in the familiar way, with simple rules.

## ■ Consider a 3 bit number:

- $001 + "-1" = 0$
- $001 + \textcolor{red}{111} = 0$ 
  - Remember  $001 + 111 = 1\ 000$ , and the leading one is lost to overflow.

## ■ $"-1" = 111$

- Yep!

# Negative Numbers

## ■ Well, if 111 is -1, what is -2?

- -1 - 1
- $111 - 001 = 110$

## ■ Does that really work?

- If it does  $-2 + 2 = 0$
- $110 + 010 = 1\ 000 = 000$

## ■ -2 + 5 should be 3, right?

- $110 + 101 = 1\ 011 = 011$

# Finding $-x$ the easy way

- Given a non-negative number in binary, e.g. 5, represented with a fixed bit width, e.g. 4
  - 0101

- Find its negative by flipping each bit and adding 1
  - 0101      This is 5
  - 1010      This is the “ones complement of 5”, e.g. 5 with bits flipped
  - 1011      This is the “twos complement of 5”, e.g. 5 with the bits flipped and 1 added
  - $0101 + 1011 = 1\ 0000 = 0000$
  - $-x = \sim x + 1$
- Because of the fixed width, the “two’s complement” of a number can be used as its negative.

# Why Does This Work?

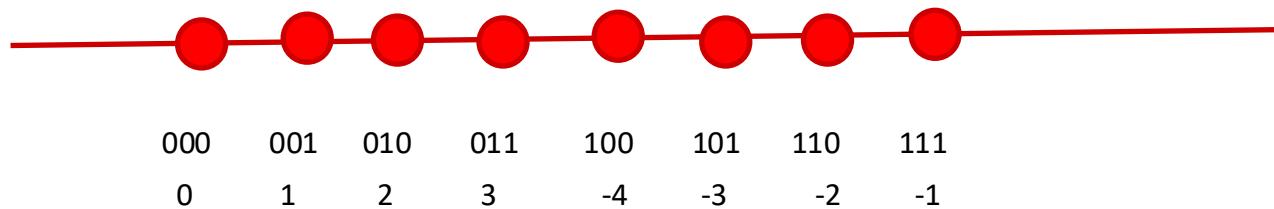
- Consider any number and its (ones) complement:
  - 0101
  - 1010
- They are called complements because complementary bits are set. As a result, if they are added, all bits are necessarily set:
  - $0101 + 1010 = 1111$
- Adding 1 to the sum of a number and its complement necessarily results in a 0 due to overflow
  - $(0101 + 1010) + 1 = 1111 + 1 = 1\ 0000 = 0000$
- And if  $x + y = 0$ ,  $y$  must equal  $-x$

# Why Does This Work? *Cont.*

- If  $x + y = 0$ 
  - $y$  must equal  $-x$
- So if  $x + (\text{Complement}(x) + 1) = 0$ 
  - $\text{Complement}(x) + 1$  must equal  $-x$
- Another way of looking at it:
  - if  $x + (\text{Complement}(x) + 1) = 0$
  - $x + \text{Complement}(x) = -1$
  - $x = -1 - \text{Complement}(x)$
  - $-x = 1 + \text{Complement}(x)$

# Visualizing Two's Complement

- Numbers “wrap around” with -1 at the very end



- A few things to note:

- All negative numbers start with a “1”
  - E.g. 100 is “-4”
- You can view the leading “1” as introducing a “-4”
  - E.g.  $101 = 1 * -4 + 0 * 2 + 1 * 1 = -3$
  - But  $010 = 0 * -4 + 1 * 2 + 0 * 1 = 2$
- -4 is missing a positive partner

# Complement & Increment Examples

$x = 0$

	Decimal	Hex	Binary
0	0	00 00	00000000 00000000
$\sim 0$	-1	FF FF	11111111 11111111
$\sim 0+1$	0	00 00	00000000 00000000

$x = T_{\min}$  (The most negative two's complement number)

	Decimal	Hex	Binary
$x$	-32768	80 00	10000000 00000000
$\sim x$	32767	7F FF	01111111 11111111
$\sim x+1$	-32768	80 00	10000000 00000000

Canonical counter example

# Encoding Integers: Dense Form

Unsigned

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

Two's Complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

```
short int x = 15213;  
short int y = -15213;
```

Sign  
Bit

- C does not mandate using two's complement
  - But, most machines do, and we will assume so
- C short (2 bytes long)

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
y	-15213	C4 93	11000100 10010011

## ■ Sign Bit

- For 2's complement, most significant bit indicates sign
  - 0 for nonnegative, 1 for negative

# Numeric Ranges

## ■ Unsigned Values

- $UMin = 0$   
000...0
- $UMax = 2^w - 1$   
111...1

## ■ Two's Complement Values

- $TMin = -2^{w-1}$   
100...0
- $TMax = 2^{w-1} - 1$   
011...1
- Minus 1  
111...1

Values for  $W = 16$

	Decimal	Hex	Binary
<b>UMax</b>	<b>65535</b>	<b>FF FF</b>	<b>11111111 11111111</b>
<b>TMax</b>	<b>32767</b>	<b>7F FF</b>	<b>01111111 11111111</b>
<b>TMin</b>	<b>-32768</b>	<b>80 00</b>	<b>10000000 00000000</b>
<b>-1</b>	<b>-1</b>	<b>FF FF</b>	<b>11111111 11111111</b>
<b>0</b>	<b>0</b>	<b>00 00</b>	<b>00000000 00000000</b>

**We're going to skip the quiz today  
Do it on your own if you want  
OR BETTER do the exercises at the end  
of the lecture**

**Check out:**

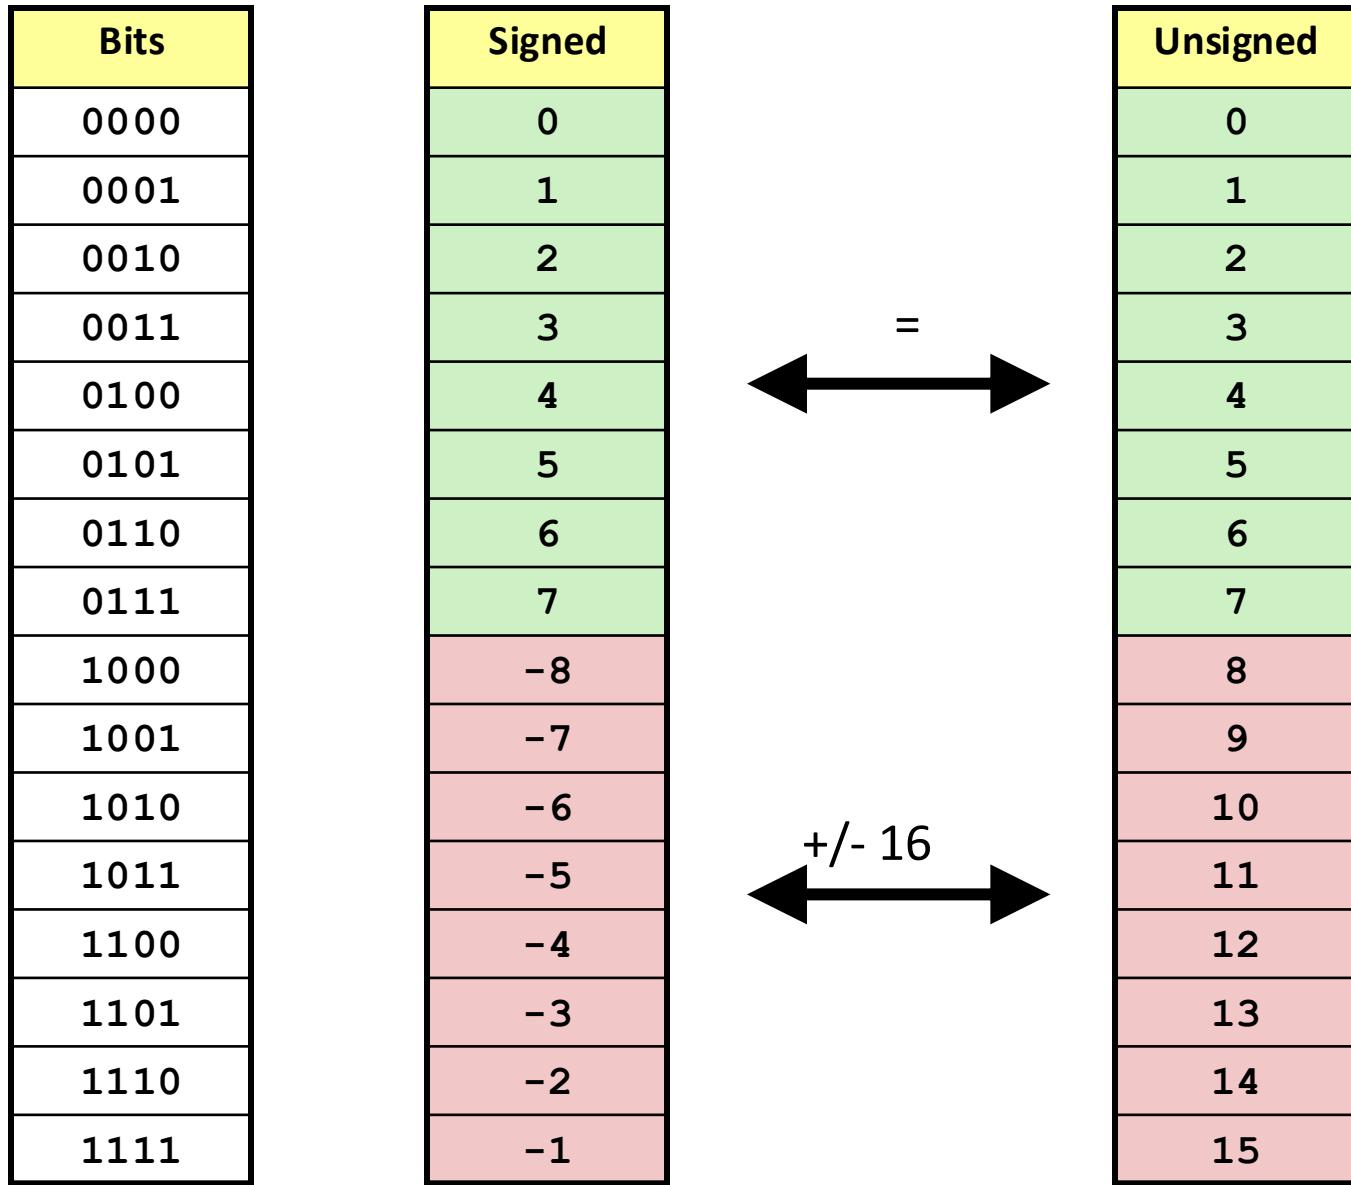
**<https://canvas.cmu.edu/courses/46319/quizzes/139185>**

```
unsigned int u = 1;
```

```
int i = (int) u;
```

```
unsigned int u = 0xffffffff;  
  
int i = (int) u;  
  
printf("%d\n", i);
```

# Mapping Signed $\leftrightarrow$ Unsigned



# Signed vs. Unsigned in C

## ■ Constants

- By default are considered to be signed integers
- Unsigned if have “U” as suffix

0U, 4294967259U

## ■ Casting

- **Explicit** casting between signed & unsigned same as U2T and T2U

```
int tx, ty;  
unsigned ux, uy;  
tx = (int) ux;  
uy = (unsigned) ty;
```

- **Implicit** casting also occurs via assignments and procedure calls

```
tx = ux;           int fun(unsigned u);  
uy = ty;           uy = fun(tx);
```

# Casting Surprises

## ■ Expression Evaluation

- If there is a mix of unsigned and signed in single expression,  
*signed values implicitly cast to unsigned*
- Including comparison operations `<`, `>`, `==`, `<=`, `>=`
- Examples for  $W = 32$ : **TMIN = -2,147,483,648** , **TMAX = 2,147,483,647**

■ Constant <sub>1</sub>	Constant <sub>2</sub>	Relation	Evaluation
0	0U	==	unsigned
-1	0	<	signed
-1	0U	>	unsigned
2147483647	-2147483647-1	>	signed
2147483647U	-2147483647-1	<	unsigned
-1	-2	>	signed
(unsigned)-1	-2	>	unsigned
2147483647	2147483648U	<	unsigned
2147483647	(int) 2147483648U	>	signed

# Summary

## Casting Signed $\leftrightarrow$ Unsigned: Basic Rules

- Bit pattern is maintained
- But reinterpreted
- Can have unexpected effects: adding or subtracting  $2^w$
  
- Expression containing signed and unsigned int
  - int is cast to unsigned!!

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - **Expanding, truncating**
  - Addition, negation, multiplication, shifting
- Byte Ordering

```
uint8_t su = 1;
```

```
uint16_t bu = su;
```

```
int8_t su = -1;
```

```
int16_t bu = su;
```

```
int8_t su = -1;
```

```
int16_t bu = su;
```

su is 0xff

bu better be 0xffff

How?

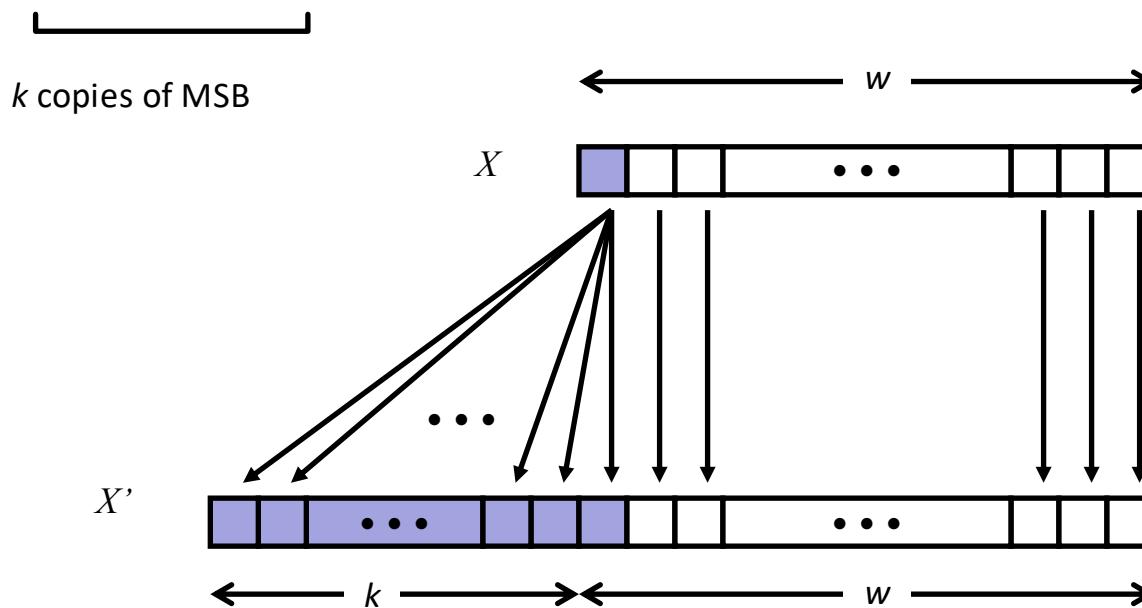
# Sign Extension

## ■ Task:

- Given  $w$ -bit signed integer  $x$
- Convert it to  $w+k$ -bit integer with same value

## ■ Rule:

- Make  $k$  copies of sign bit:
- $X' = x_{w-1}, \dots, x_{w-1}, x_{w-1}, x_{w-2}, \dots, x_0$



```
uint16_t bu = 1;
```

```
uint8_t su = bu;
```

```
uint16_t bu = 1000;
```

```
uint8_t su = bu; ???
```

bu = 0x8000 (0b100..)

= negative max

su = bu ???

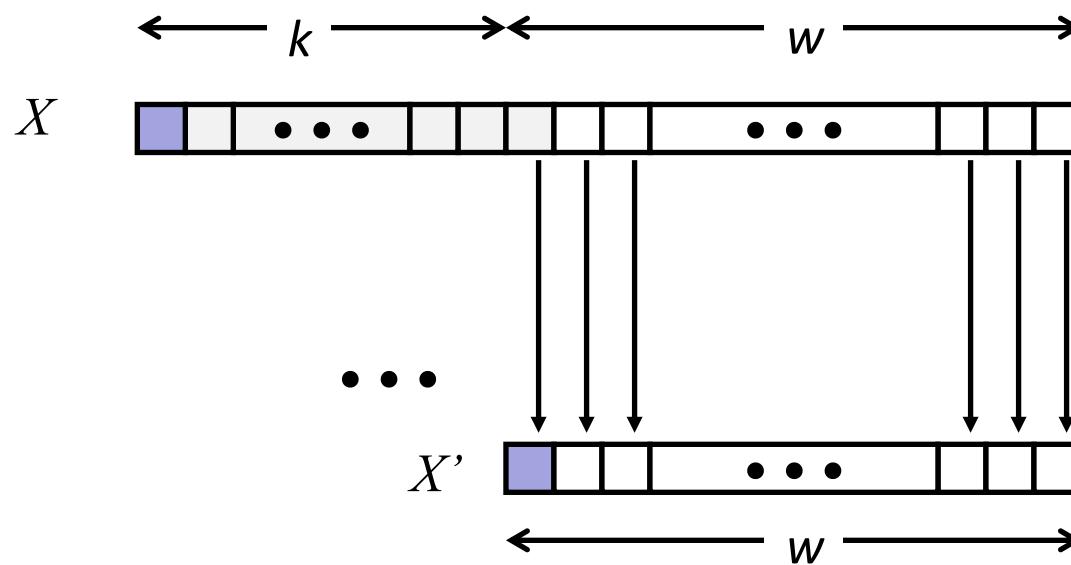
# Truncation

## ■ Task:

- Given  $k+w$ -bit signed or unsigned integer  $X$
- Convert it to  $w$ -bit integer  $X'$  with same value for “small enough”  $X$

## ■ Rule:

- Drop top  $k$  bits:
- $X' = x_{w-1}, x_{w-2}, \dots, x_0$



# Truncation: Simple Example

No sign change

	-16	8	4	2	1
2 =	0	0	0	1	0

	-8	4	2	1
2 =	0	0	1	0

$$2 \bmod 16 = 2$$

	-16	8	4	2	1
-6 =	1	1	0	1	0

	-8	4	2	1
-6 =	1	0	1	0

$$-6 \bmod 16 = 26U \bmod 16 = 10U = -6$$

Sign change

	-16	8	4	2	1
10 =	0	1	0	1	0

	-8	4	2	1
-6 =	1	0	1	0

$$10 \bmod 16 = 10U \bmod 16 = 10U = -6$$

	-16	8	4	2	1
-10 =	1	0	1	1	0

	-8	4	2	1
6 =	0	1	1	0

$$-10 \bmod 16 = 22U \bmod 16 = 6U = 6$$

# **Summary:**

## **Expanding, Truncating: Basic Rules**

### **■ Expanding (e.g., short int to int)**

- Unsigned: zeros added
- Signed: sign extension
- Both yield expected result

### **■ Truncating (e.g., unsigned to signed short)**

- Unsigned/signed: bits are truncated
- Result reinterpreted
- Unsigned: mod operation
- Signed: similar to mod
- For small (in magnitude) numbers yields expected behavior

# Today: Bits, Bytes, and Integers

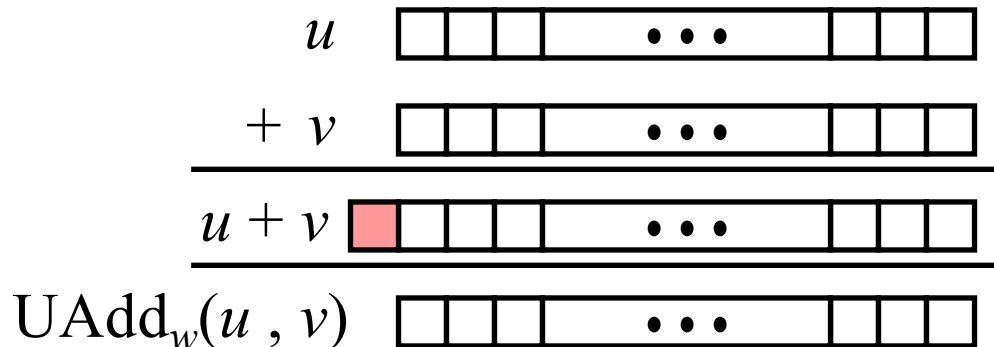
- Representing information as bits
- Bit-level manipulations
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - **Addition, negation, multiplication, shifting**
- Byte Ordering

# Unsigned Addition

Operands:  $w$  bits

True Sum:  $w+1$  bits

Discard Carry:  $w$  bits



## ■ Standard Addition Function

- Ignores carry output

## ■ Implements Modular Arithmetic

$$s = \text{UAdd}_w(u, v) = u + v \bmod 2^w$$

unsigned char	1110 1001	E9	223
	+ 1101 0101	+ D5	+ 213
	<hr/>	<hr/>	<hr/>
	1 1011 1110	1BE	446
	<hr/>	<hr/>	<hr/>
	1011 1110	BE	190

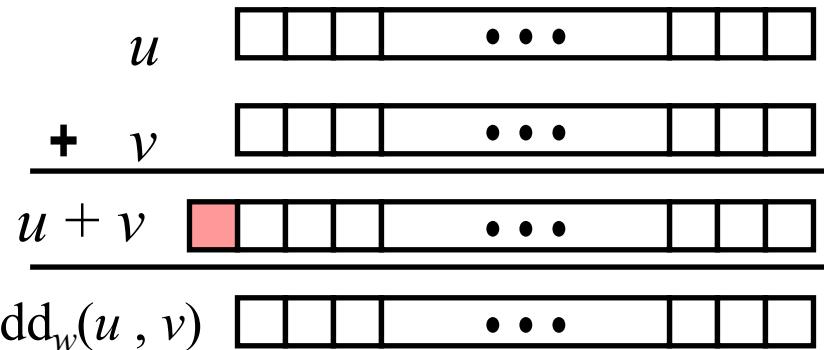
Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

# Two's Complement Addition

Operands:  $w$  bits

True Sum:  $w+1$  bits

Discard Carry:  $w$  bits



## ■ TAdd and UAdd have Identical Bit-Level Behavior

- Signed vs. unsigned addition in C:

```
int s, t, u, v;  
s = (int) ((unsigned) u + (unsigned) v);  
t = u + v
```

- Will give  $s == t$

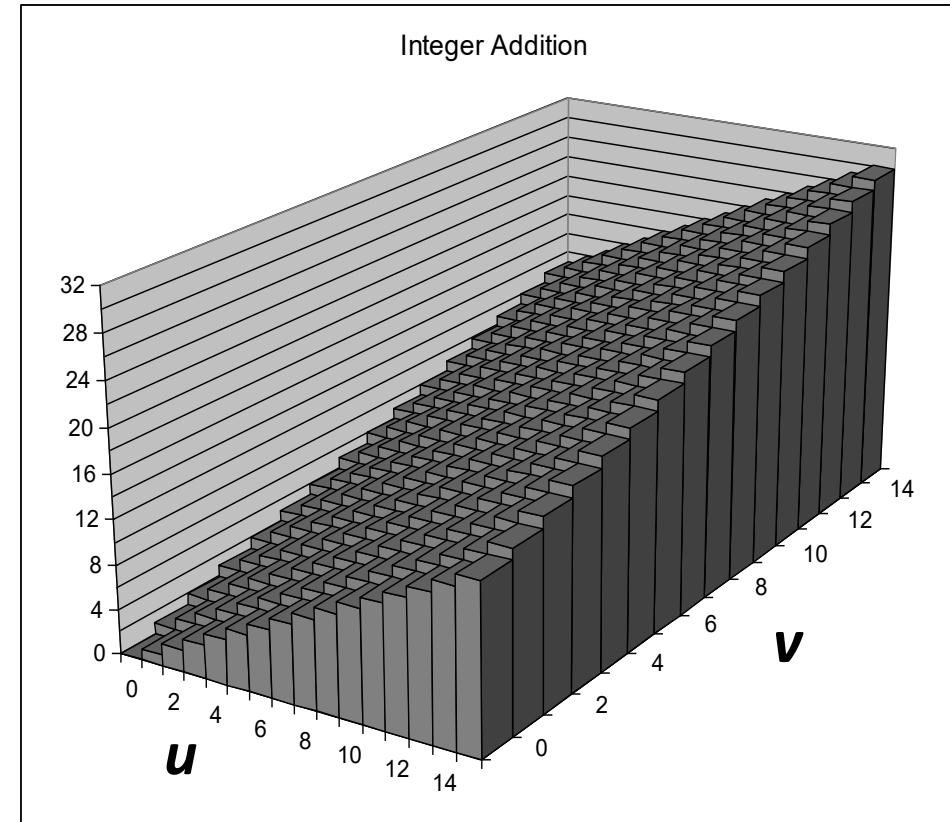
$$\begin{array}{r} 1110 \ 1001 \\ + \ 1101 \ 0101 \\ \hline \textcolor{red}{1} \ \textcolor{red}{1011} \ \textcolor{red}{1110} \end{array} \qquad \begin{array}{r} \text{E9} \\ + \ D5 \\ \hline \textcolor{red}{1BE} \end{array} \qquad \begin{array}{r} -23 \\ + \ -43 \\ \hline \textcolor{red}{-66} \end{array}$$
$$\begin{array}{r} 1011 \ 1110 \\ \hline \textcolor{red}{BE} \end{array} \qquad \begin{array}{r} \\ \\ \end{array} \qquad \begin{array}{r} \\ \\ \end{array}$$

# Visualizing “True Sum” Integer Addition

## ■ Integer Addition

- 4-bit integers  $u, v$
- Compute true sum  
 $\text{Add}_4(u, v)$
- Values increase linearly  
with  $u$  and  $v$
- Forms planar surface

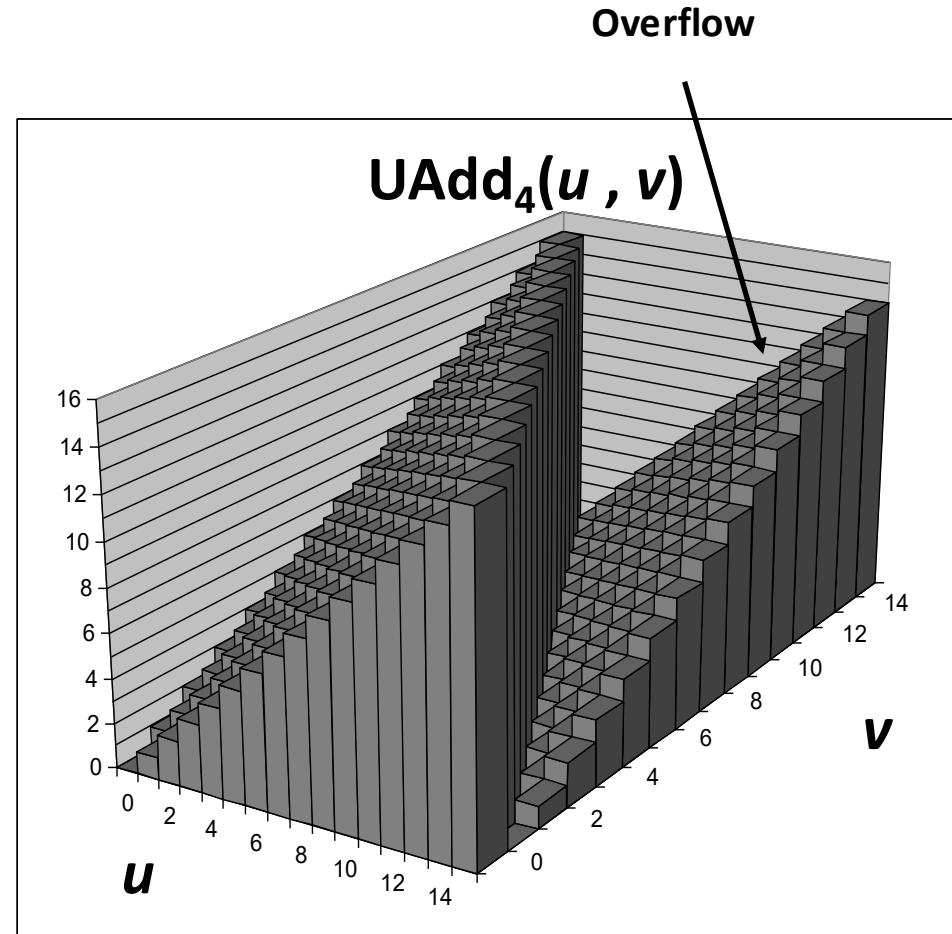
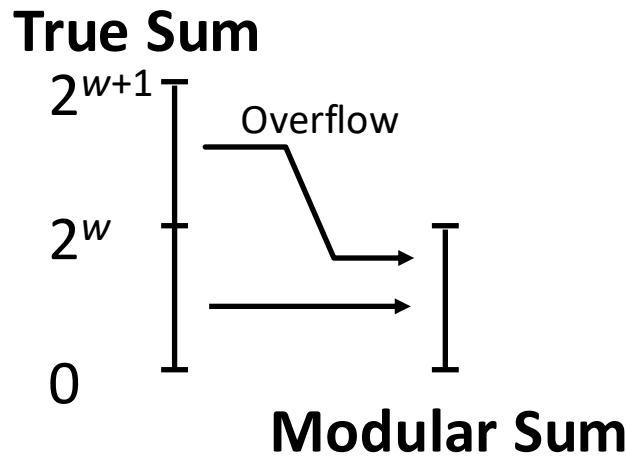
$\text{Add}_4(u, v)$



# Visualizing Unsigned Addition

## Wraps Around

- If true sum  $\geq 2^w$
- At most once



# Visualizing 2's Complement Addition

## ■ Values

- 4-bit two's comp.
- Range from -8 to +7

## ■ Wraps Around

- If sum  $\geq 2^{w-1}$ 
  - Becomes negative
  - At most once
- If sum  $< -2^{w-1}$ 
  - Becomes positive
  - At most once

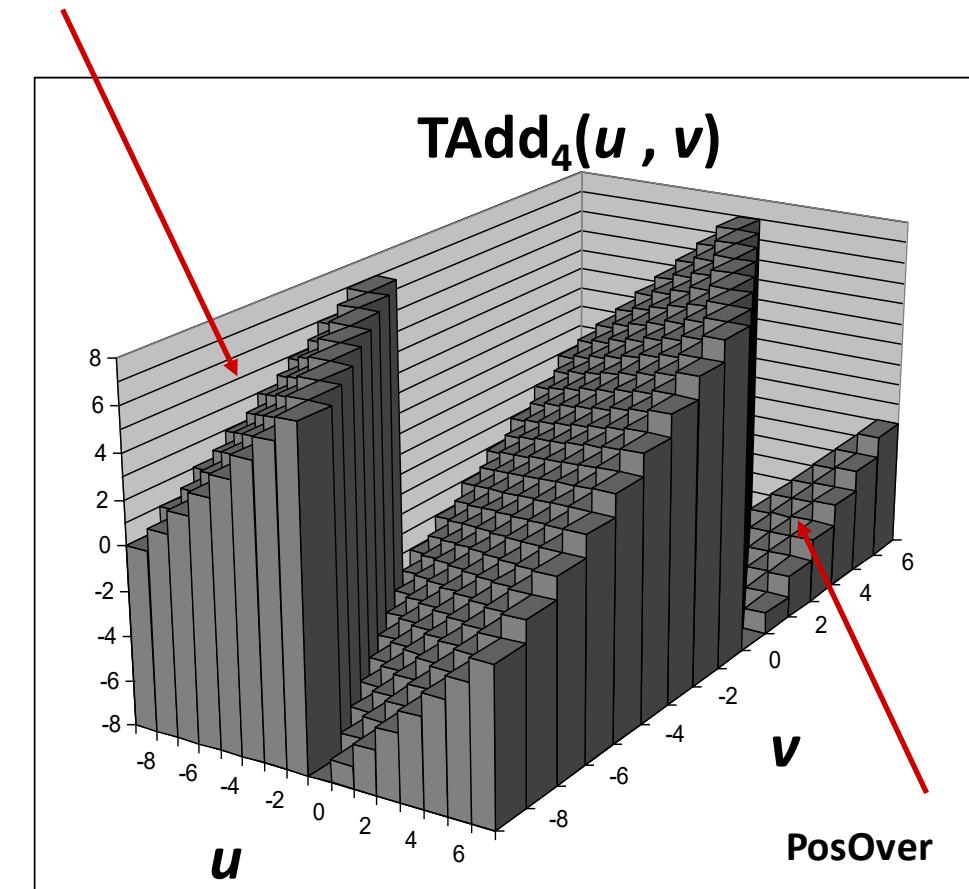
NegOver

TAdd<sub>4</sub>( $u$ ,  $v$ )

$u$

PosOver

$v$



# Multiplication

- Goal: Computing Product of  $w$ -bit numbers  $x, y$ 
  - Either signed or unsigned
- Result: Same as computing ideal, exact result  $x^*y$  and keeping  $w$  lower bits.
- Ideal,exact results can be bigger than  $w$  bits
  - Worst case is up to  $2w$  bits
    - Unsigned, because all bits are magnitude
    - Signed, but only for  $T_{min}^*T_{min}$ , because anything added to  $T_{min}$  reduces its magnitude and  $T_{max}$  is less than  $T_{min}$ .
- So, maintaining exact results...
  - would need to keep expanding word size with each product computed
  - Impossible in hardware (at least without limits), as all resources are finite
  - In practice, is done in software, if needed
    - e.g., by “arbitrary precision” arithmetic packages

# Power-of-2 Multiply with Shift

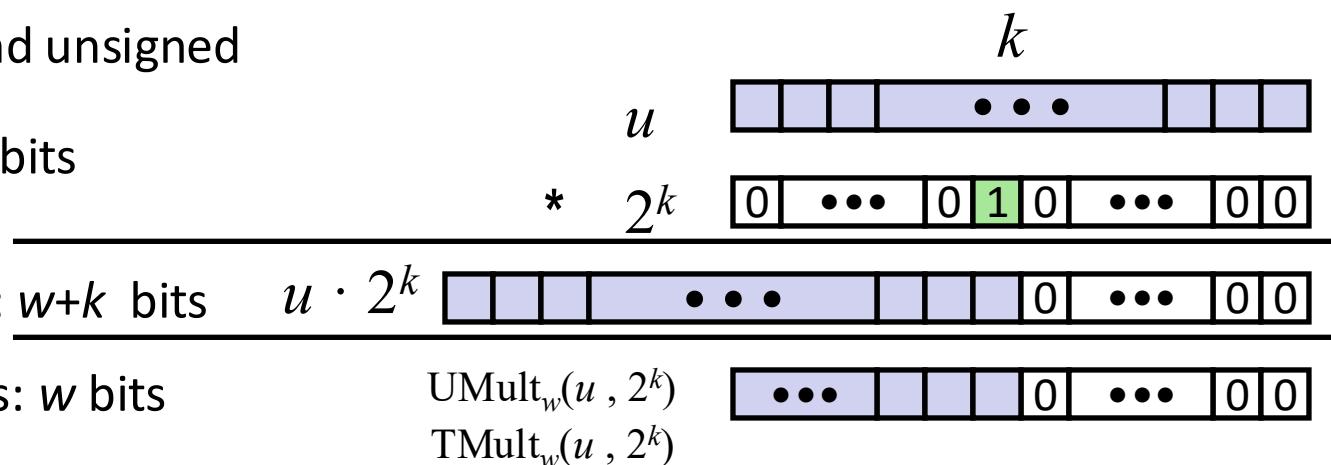
## ■ Operation

- $u \ll k$  gives  $u * 2^k$
- Both signed and unsigned

Operands:  $w$  bits

True Product:  $w+k$  bits

Discard  $k$  bits:  $w$  bits



## ■ Examples

- $u \ll 3 == u * 8$
- $(u \ll 5) - (u \ll 3) == u * 24$
- Most machines shift and add faster than multiply
  - Compiler generates this code automatically

# Dividing by shifting

$$4 / 2 = 2$$

$$0100 \gg 1 = 0010$$

$$4 / 4 = 1$$

$$0100 \gg 2 = 0001$$

$$N / 2^k$$

$$N \gg k$$

# How does this round?

$$5 / 2 = 2$$

$$0101 \gg 1 = 0010$$

$$-5 / 2 = ?$$

$$\begin{aligned} 1011 \gg 1 &= 1101 \\ &= -3 \end{aligned}$$

It rounds *Lower*

# How do you round to zero?

Add a bias!

$$\begin{array}{r} -5 \\ + 1 \end{array}$$

$$\begin{array}{r} 1011 \\ + 1 \end{array} = 1100 = -4$$

$$\begin{array}{r} 1100 \\ \gg 1 \end{array} = 1110 = -2$$

$$\begin{array}{r} -4 + 1 \% 2 \end{array}$$

$$\begin{array}{r} 1100 + 1 \end{array} = 1101$$

$$\begin{array}{r} 1101 \gg 1 \end{array} = 1110 = -2$$

# Generalized bias for dividing negative #s

$$\lfloor u / 2^k \rfloor$$

=

$$(u + (1 \ll k) - 1) \gg k$$

Rounds to zero

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
- Byte Ordering

# Byte Ordering

- So, how are the bytes within a multi-byte word ordered in memory?

- Conventions

- Big Endian: Sun (Oracle SPARC), PPC Mac, *Internet*
  - Least significant byte has highest address
- Little Endian: *x86*, ARM processors running Android, iOS, and Linux
  - Least significant byte has lowest address

- Becomes a concern when data is communicated

- Over a network, via files, etc.

- Important notes

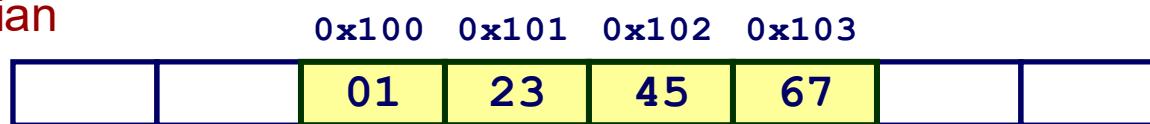
- Bits are not reversed, as the low order bit is the reference point.
- Doesn't affect chars, or strings (arrays of chars), as chars are only one byte

# Byte Ordering Example

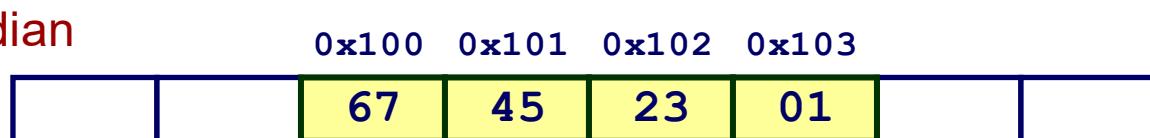
## ■ Example

- Variable x has 4-byte value of 0x01234567
- Address given by &x is 0x100

Big Endian



Little Endian



# Reading Byte-Reversed Listings

## ■ Disassembly

- Text representation of binary machine code
- Generated by program that reads the machine code

## ■ Example Fragment

Address	Instruction Code	Assembly Rendition
8048365:	5b	pop %ebx
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx
804836c:	83 bb 28 00 00 00	cmpl \$0x0,0x28(%ebx)

## ■ Deciphering Numbers

- Value:
- Pad to 32 bits:
- Split into bytes:
- Reverse:

0x12ab

0x000012ab

00 00 12 ab

ab 12 00 00

# Today: Bits, Bytes, and Integers

- Representing information as bits CSAPP 2.1
- Bit-level manipulations
- Integers CSAPP 2.2
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting CSAPP 2.3
- Byte Ordering CSAPP 2.1.3

Some learning exercises:

<https://wiredream.com/learning-to-bit-twiddle/>