

# **Machine-Level Programming I**

15-213/15-513: Introduction to Computer Systems

3rd Lecture, September 2, 2025

**in-person,  
graded quiz**

**at the start of class Thursday**

$$(\sim a + 1) = -a$$

**movq %rax, 0x20(%rdi, %rsi, 2)**

# Today: Machine Programming I: Basics

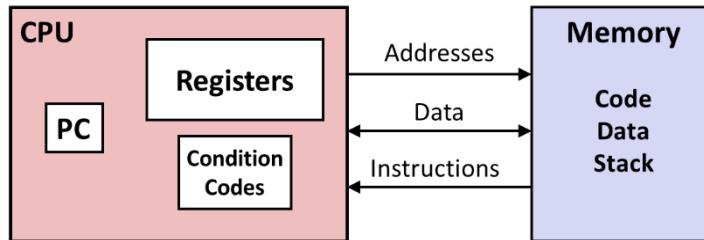
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- Condition codes and jumps
- C, assembly, machine code

# Context: Levels of Abstraction

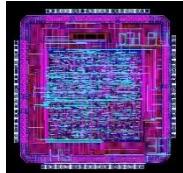
C programmer

```
#include <stdio.h>
int main(){
    int i, n = 10, t1 = 0, t2 = 1, nxt;
    for (i = 1; i <= n; ++i){
        printf("%d, ", t1);
        nxt = t1 + t2;
        t1 = t2;
        t2 = nxt; }
    return 0; }
```

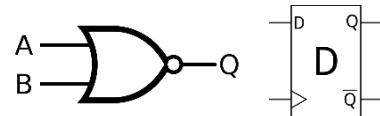
Assembly programmer



Computer Designer



Gates, clocks, circuit layout, ...



# CPUs store data in Registers and Memory

- Registers: Super-fast memory right on the CPU core itself.
  - About 16 named registers on x86
- Many instructions operate on registers
- Registers are fixed-width, often 32 or 64 bits
- “int x = 5; x \*= y; x++” – x is probably stored in a *register* while being operated on frequently.
- Registers are accessed by their name
- “Memory” means DRAM, typically off-chip storage, bigger and cheaper and slower. When you call “malloc”, you usually get this.
- Memory is accessed by (byte) address

# Introduction to Assembly

- **Assembly is a programming language**
  - It has types
  - It has control flow shape
  - It is also an intermediate, very limited verbs, need to be explicit
- **Assembly is defined by the ISA  
(instruction set architecture)**
  - x86, ARM, RISC-V, etc
- **Implemented by the microarchitecture**
  - e.g., a “core i7” x86 processor

# Beginning Grammar

- Say something in assembly:

Operation (op)      Register names

The diagram illustrates the breakdown of the assembly instruction 'add %rbx, %rax'. It features two labels at the top: 'Operation (op)' pointing to the first part of the instruction, and 'Register names' pointing to the second and third parts. Below these labels is a yellow rectangular box containing the tokens 'add', '%rbx,' and '%rax'. To the right of this box, the word 'is' is centered, followed by the assembly instruction 'rax += rbx'.

add      %rbx,      %rax

is

rax += rbx

# Nouns

## ■ Registers

- Special parts of the CPU to hold small amounts of data, like local variables
- Names are very specific:

%r?x (a,b,c,d)

%r?i (d,s)

%r?p (b,i,s) // Special meanings

%r? (8-15)

## ■ Memory

- Everything else; accessed by *address*.

# X86 registers – some history

Intel created x86 in the 16-bit microprocessor 8086.

Its registers held 16 bits. So:

“ax”, the “accumulator”, stored 16 bits of data.

Initially the registers had a defined purpose. (They mostly don't now.)

In 1985, Intel introduced the 32-bit 80386, and “e”xtended the registers to 32 bits, creating “eax”

In 2003, AMD introduced x86-64, growing registers to 64 bits and creating “rax”

# x86-64 Integer Registers (**reference**)

%rax	%eax
------	------

%rbx	%ebx
------	------

%rcx	%ecx
------	------

%rdx	%edx
------	------

%rsi	%esi
------	------

%rdi	%edi
------	------

%rsp	%esp
------	------

%rbp	%ebp
------	------

%r8	%r8d
-----	------

%r9	%r9d
-----	------

%r10	%r10d
------	-------

%r11	%r11d
------	-------

%r12	%r12d
------	-------

%r13	%r13d
------	-------

%r14	%r14d
------	-------

%r15	%r15d
------	-------

- Can reference low-order 4 bytes (also low-order 1 & 2 bytes)
- Not part of memory (or cache)

# Verbs (i.e., operations)

- Transfer data between memory and register
  - Load data from memory into register
  - Store register data into memory
- Perform arithmetic function on register or memory data
- Transfer control
  - Unconditional jumps to/from procedures
  - Conditional branches
  - Indirect branches

# Verbs have optional suffix

- Operation suffix specifies the size(s) involved
  - There are also different register names for the different sizes

Prefix	Register Name	Size (bytes)
q	r..	8
l	e.. (usually)	4
w	see reference	2
b	see reference	1

# Moving Data

- $x = 5$

**mov \$5, %rax**

- $x = y$

**mov %rcx, %rdx**

- $x = *p$

**mov (%rsi), %r8**

# Memory

- Parentheses always denote a memory address computation

(%rdi) -> \*ptr

- Most operations will then access memory

# `movq` Operand Combinations (**reference**)

	Source	Dest	Src,Dest	C Analog
<code>movq</code>	<i>Imm</i>	<i>Reg</i>	<code>movq \$0x4,%rax</code>	<code>temp = 0x4;</code>
	<i>Imm</i>	<i>Mem</i>	<code>movq \$-147,(%rax)</code>	<code>*p = -147;</code>
	<i>Reg</i>	<i>Reg</i>	<code>movq %rax,%rdx</code>	<code>temp2 = temp1;</code>
	<i>Reg</i>	<i>Mem</i>	<code>movq %rax,(%rdx)</code>	<code>*p = temp;</code>
	<i>Mem</i>	<i>Reg</i>	<code>movq (%rax),%rdx</code>	<code>temp = *p;</code>

*Cannot do memory-memory transfer with a single instruction*

# Reading Assembly Example

```
void  
whatAmI(<type> a, <type> b)  
{  
   ????  
}
```

%rdi

%rsi

whatAmI:

```
    movq    (%rdi), %rax  
    movq    (%rsi), %rdx  
    movq    %rdx, (%rdi)  
    movq    %rax, (%rsi)  
    ret
```

What is <type>?

Find %rdi in the assembly

How is it used?

# Reading Assembly Example

```
void swap
    (long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

**swap:**

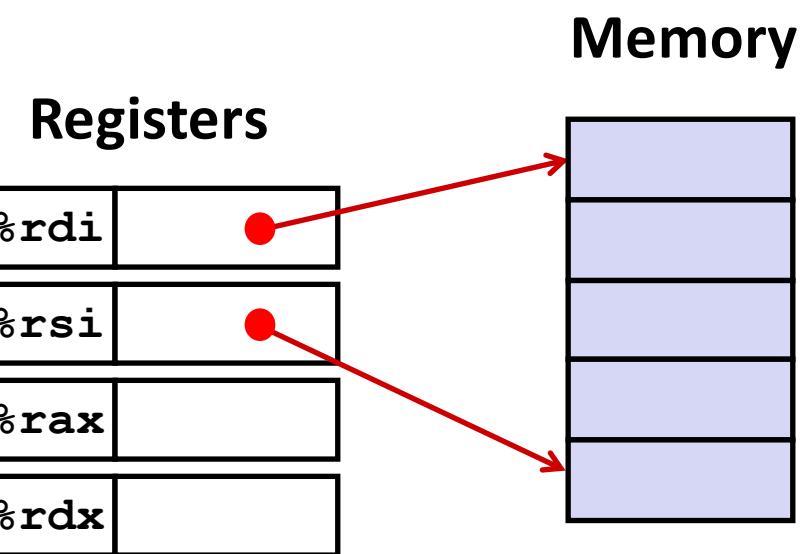
movq	(%rdi), %rax
movq	(%rsi), %rdx
movq	%rdx, (%rdi)
movq	%rax, (%rsi)
ret	

We have now written the “pseudo-C” for the assembly, using the types and replacing each register with a “variable name”.

# Understanding Swap()

```
void swap
    (long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

Register	Value
%rdi	xp
%rsi	yp
%rax	t0
%rdx	t1



swap:

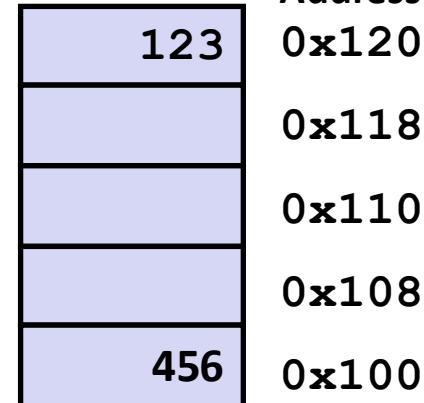
```
        movq    (%rdi), %rax    # t0 = *xp
        movq    (%rsi), %rdx    # t1 = *yp
        movq    %rdx, (%rdi)    # *xp = t1
        movq    %rax, (%rsi)    # *yp = t0
        ret
```

# Understanding Swap()

Registers

%rdi	0x120
%rsi	0x100
%rax	
%rdx	

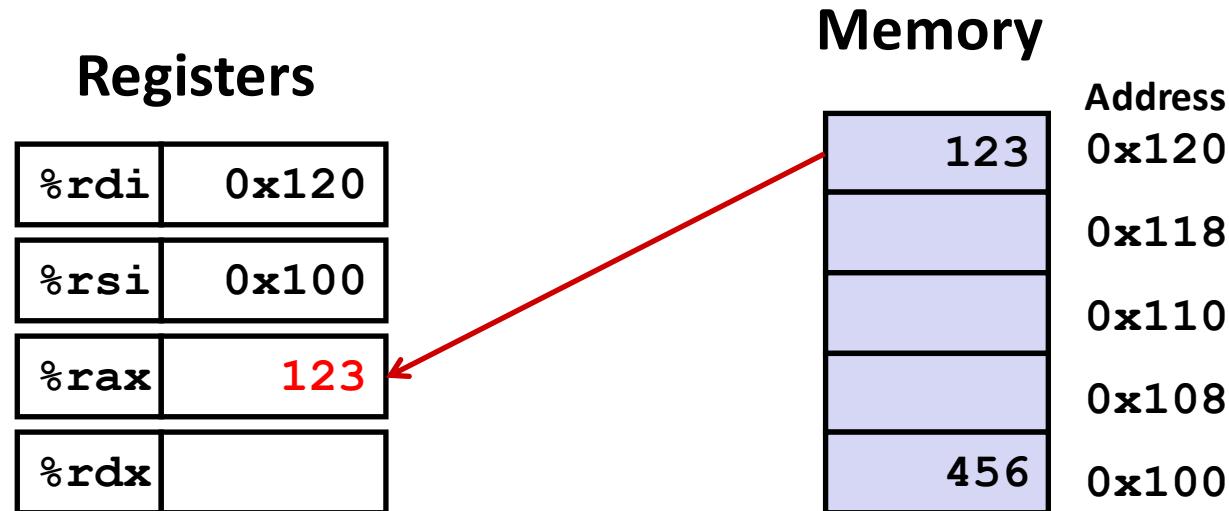
Memory



swap:

```
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
```

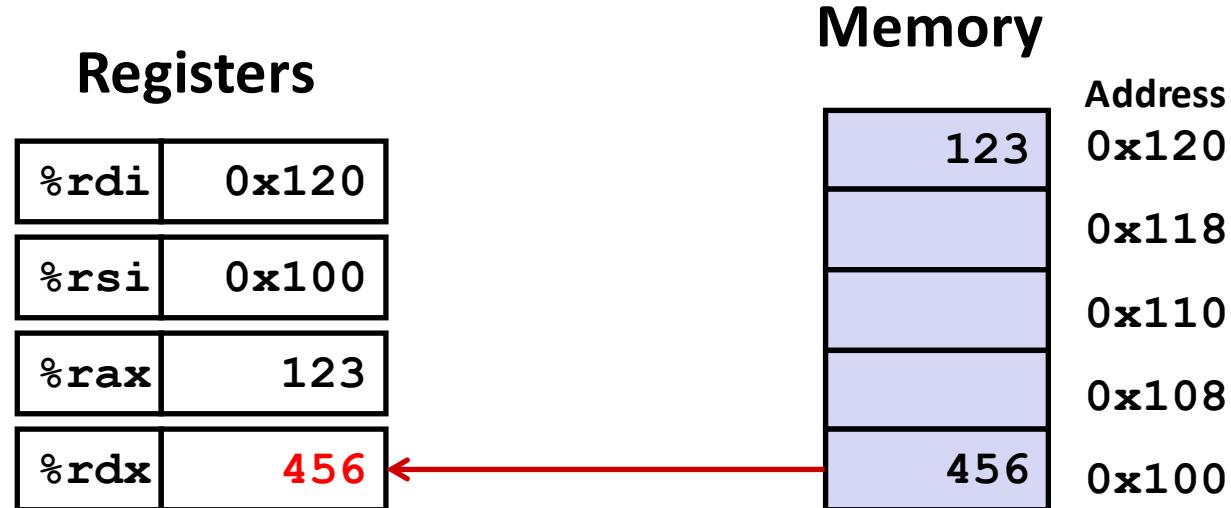
# Understanding Swap()



swap:

```
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
```

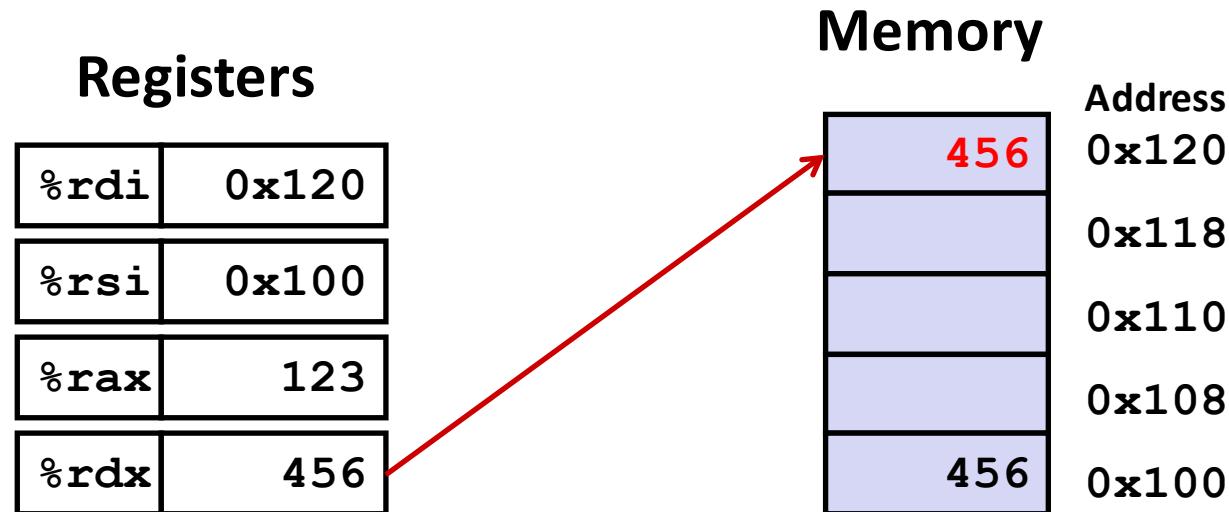
# Understanding Swap()



swap:

```
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
```

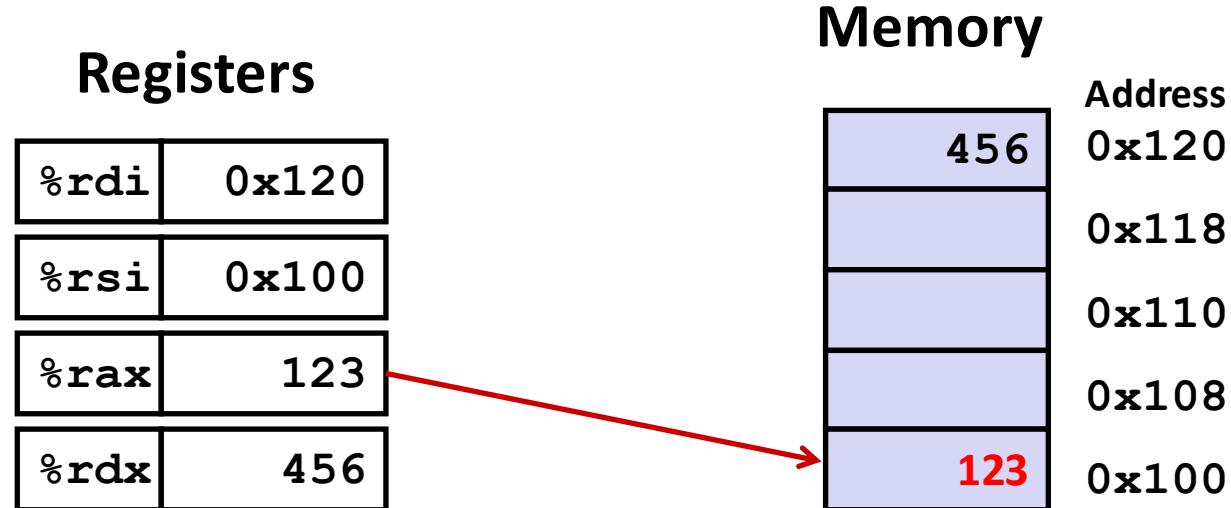
# Understanding Swap()



swap:

```
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
```

# Understanding Swap()



swap:

```
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
```

# More memory grammar

```
movq    (%rdi) , %rax  
x = *rdi
```

```
movq    (%rdi, %rsi) , %rax  
x = *(rdi + rsi)
```

```
movq    (%rdi, %rsi, 4) , %rax  
x = *(rdi + 4 * rsi)
```

```
movq    0x80(%rdi, %rsi, 4) , %rax  
x = *(rdi + 4 * rsi + 0x80)
```

# The memory ops make arrays easy

```
char *d;  
int y;  
char a = '\0';  
d[y] = a;
```

```
movq    %rax, (%rdi, %rsi),  
%rax
```

# When array type is larger than a byte

```
int *d;  
int s;  
char a = 67;  
d[y] = a;
```

```
movq    (%rdi, %rsi, 4), %rax  
  
sizeof(int) == 4
```

# Generalized memory grammar

**D(Rb,Ri,S)**

**Mem[Reg[Rb]+S\*Reg[Ri]+ D]**

**D** Constant “displacement” 1, 2, or 4 bytes

**Rb** Base register: Any of 16 integer registers

**Ri** Index register: Any, except for %**rsp**

**S** Scale: 1, 2, 4, or 8

**If a value is “missing”, it is the identity element**

- +0 or \*1

# Address Computation Examples

<code>%rdx</code>	<code>0xf000</code>
<code>%rcx</code>	<code>0x0100</code>

**D(Rb,Ri,S)**

**Mem[Reg[Rb]+S\*Reg[Ri]+ D]**

- D: Constant “displacement” 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for `%rsp`
- S: Scale: 1, 2, 4, or 8 (*why these numbers?*)

Expression	Address Computation	Address
<code>0x8(%rdx)</code>		
<code>(%rdx,%rcx)</code>		
<code>(%rdx,%rcx,4)</code>		
<code>0x80(,%rdx,2)</code>		

# More Verbs

## ■ Two operands (think datalab)

<i>Format</i>	<i>Computation</i>	
addq <i>Src,Dest</i>	$Dest = Dest + Src$	
salq <i>Src,Dest</i>	$Dest = Dest << Src$	<i>Also called shlq</i>
sarq <i>Src,Dest</i>	$Dest = Dest >> Src$	<i>Arithmetic</i>
shrq <i>Src,Dest</i>	$Dest = Dest >> Src$	<i>Logical</i>
...		

## ■ One operand (also datalab)

incq <i>Dest</i>	$Dest = Dest + 1$
decq <i>Dest</i>	$Dest = Dest - 1$
negq <i>Dest</i>	$Dest = - Dest$
notq <i>Dest</i>	$Dest = \sim Dest$

# Address Computation Instruction

## ■ **leaq *Src*, *Dst***

- *Src* is address mode expression
- Set *Dst* to address denoted by expression

## ■ **Uses**

- Computing addresses without a memory reference
  - E.g., translation of `p = &x[i];`
- Computing arithmetic expressions of the form  $x + k*y$ 
  - $k = 1, 2, 4, \text{ or } 8$

## ■ **Example**

```
long m12(long x)
{
    return x*12;
}
```

Converted to ASM by compiler:

```
leaq (%rdi,%rdi,2), %rax # t = x+2*x
salq $2, %rax           # return t<<2
```

# Registers have conventions

## ■ Architectural requirements:

- %rip – program counter / instruction pointer
  - Only points to the next instruction to execute
- %rsp – stack pointer
  - %rbp – sometimes also stack related

## ■ “Calling Convention” (must be agreed upon for interop!)

- %rax – return value
- %rdi – first argument
- %rsi – second argument
- %rdx – third argument
- ...

# Putting it Together

## ■ Arguments?

- How many?
- What type(s)?

## ■ Return value?

## ■ Local variables?

**arith:**

<b>leaq</b>	(%rdi,%rsi), %rax
<b>addq</b>	%rdx, %rax
<b>leaq</b>	(%rsi,%rsi,2), %rdx
<b>salq</b>	\$4, %rdx
<b>leaq</b>	4(%rdi,%rdx), %rcx
<b>imulq</b>	%rcx, %rax
<b>ret</b>	

## Interesting Instructions

- **leaq**: address computation
- **salq**: shift
- **imulq**: multiplication
  - But, only used once

# Arithmetic Expression Example

```
long arith  
(long x, long y, long z)  
{  
    long t1 = x+y;  
    long t2 = z+t1;  
    long t3 = x+4;  
    long t4 = y * 48;  
    long t5 = t3 + t4;  
    long rval = t2 * t5;  
    return rval;  
}
```

```
arith:  
    leaq    (%rdi,%rsi), %rax  
    addq    %rdx, %rax  
    leaq    (%rsi,%rsi,2), %rdx  
    salq    $4, %rdx  
    leaq    4(%rdi,%rdx), %rcx  
    imulq   %rcx, %rax  
    ret
```

## Interesting Instructions

- **leaq**: address computation
- **salq**: shift
- **imulq**: multiplication
  - But, only used once

# Understanding Arithmetic Expression Example

```
long arith  
(long x, long y, long z)  
{  
    long t1 = x+y;  
    long t2 = z+t1;  
    long t3 = x+4;  
    long t4 = y * 48;  
    long t5 = t3 + t4;  
    long rval = t2 * t5;  
    return rval;  
}
```

```
arith:  
    leaq    (%rdi,%rsi), %rax    # t1  
    addq    %rdx, %rax          # t2  
    leaq    (%rsi,%rsi,2), %rdx  
    salq    $4, %rdx            # t4  
    leaq    4(%rdi,%rdx), %rcx  # t5  
    imulq   %rcx, %rax          # rval  
    ret
```

Register	Use(s)
%rdi	Argument x
%rsi	Argument y
%rdx	Argument z , t4
%rax	t1, t2, rval
%rcx	t5

# Expressing Inequality

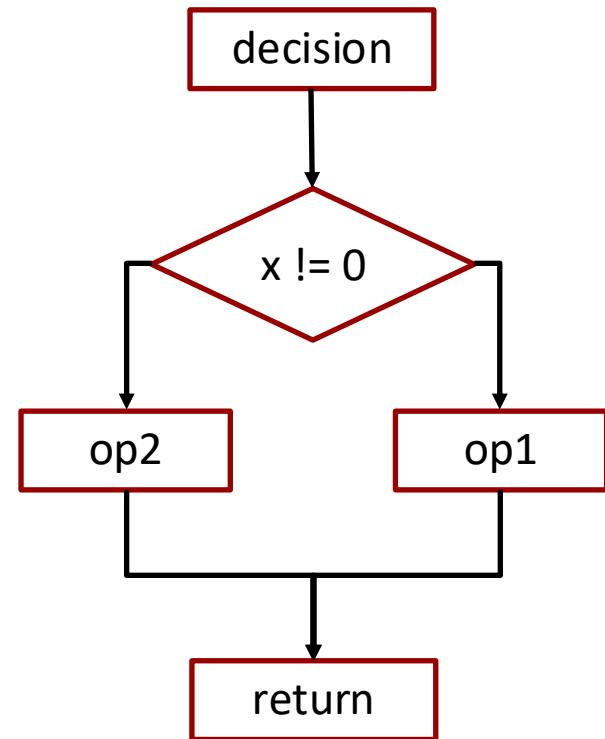
- We want to express numeric relations in assembly
  - Equals, not equals
  - Greater, less than

# Control Flow

- We also need assembly to make decisions based on these (in)equalities
  - We call the sequence of instructions executed, the control flow

# Control flow

```
extern void op1(void);  
extern void op2(void);  
  
void decision(int x) {  
    if (x) {  
        op1();  
    } else {  
        op2();  
    }  
}
```



# Control flow in assembly language

```
extern void op1(void);  
extern void op2(void);  
  
void decision(int x) {  
    if (x) {  
        op1();  
    } else {  
        op2();  
    }  
}
```

```
decision:  
    testl    %edi, %edi  
    je       .L2  
    call    op1  
    jmp     .L1  
  
.L2:  
    call    op2  
  
.L1:  
    ret
```



# Conditional “Goto”

- How did the example work?
- The first jump instruction is conditional
  - It uses processor state set by the **test** instruction
  - Processor has special registers to hold specific state
- We call this specific state, “condition codes”

# Expressing Inequality

- x86 has two instructions to set the specific state
  - cmp
  - test
- Instructions other than lea also implicitly set the state on x86

# Compare Instruction

## ■ **cmp a, b**

- Computes  $b - a$  (just like **sub**)
- Sets condition codes based on result, but...
- **Does not change b**
  
- Used for **if (a < b) { ... }**  
whenever  $b - a$  isn't needed for anything else

# Test Instruction

## ■ **test a, b**

- Computes  $b \& a$  (just like **and**)
- Sets condition codes (only SF and ZF) based on result, but...
- **Does not change b**
- Most common use: `test %rX, %rX`  
to compare `%rX` to zero
- Second most common use: `test %rX, %rY`  
tests if any of the 1-bits in `%rY` are also 1 in `%rX` (or vice versa)

# Jumping

## ■ **jX Instructions**

- Jump to different part of code depending on condition codes
- jmp – unconditional
  
- Z
  - ZERO
  
- GE
  - Greater than or equal

# Jumping (reference)

## ■ jX Instructions

- Jump to different part of code depending on condition codes

jX	Condition	Description
jmp	1	Unconditional
je	ZF	Equal / Zero
jne	~ZF	Not Equal / Not Zero
js	SF	Negative
jns	~SF	Nonnegative
jg	~(SF^OF) & ~ZF	Greater (Signed)
jge	~(SF^OF)	Greater or Equal (Signed)
jl	(SF^OF)	Less (Signed)
jle	(SF^OF)   ZF	Less or Equal (Signed)
ja	~CF & ~ZF	Above (unsigned)
jb	CF	Below (unsigned)

# Returning Condition Codes

## ■ setX will

- Set low-order byte of destination to 0 or 1 based on *combinations* of condition codes
- Does not alter remaining 7 bytes

```
int gt (long x, long y)
{
    return x > y;
}
```

```
cmpq    %rsi, %rdi      # Compare x:y
setg    %al               # Set when >
movzbl  %al, %eax       # Zero rest of %rax
ret
```

# **movXYZ – Moving to larger bit widths**

## ■ **mov with three suffixes will move to larger bit widths**

- X – {s,z}
  - Sign extend
  - Zero extend
- Y – source bit width
- Z – destination bit width

## ■ **movzbl %al, %eax**

- z – zero
- b – source is 1 byte
- l – destination is 4 bytes

# Summary

## ■ Nouns

- Registers

## ■ Verbs

- Instructions or operations

## ■ Suffixes and Annotations

- Specify the size (usually optional)
- Memory addressing mode

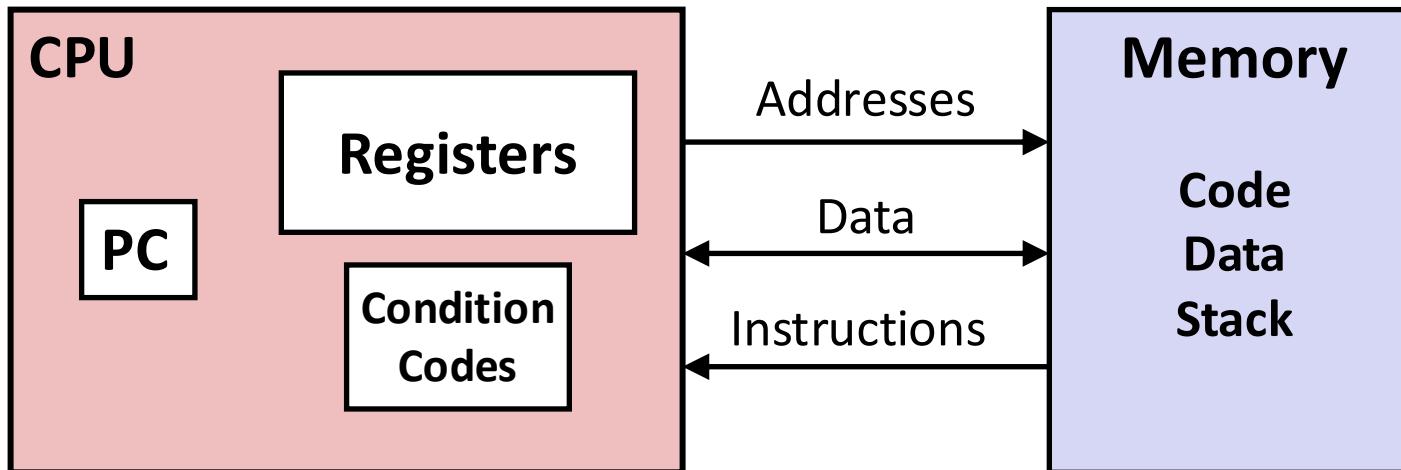
## ■ Relations

- Relies on condition codes and conditional jump (i.e., goto)

# Definitions

- **Architecture:** (also ISA: instruction set architecture) The parts of a processor design that one needs to understand for writing assembly/machine code.
  - Examples: instruction set specification, registers
- **Microarchitecture:** Implementation of the architecture
  - Examples: cache sizes and core frequency
- **Code Forms:**
  - **Machine Code:** The byte-level programs that a processor executes
  - **Assembly Code:** A text representation of machine code
- **Example ISAs:**
  - Intel: x86, IA32, Itanium, x86-64
  - ARM: Used in almost all mobile phones
  - RISC V: New open-source ISA

# Assembly/Machine Code View



## Programmer-Visible State

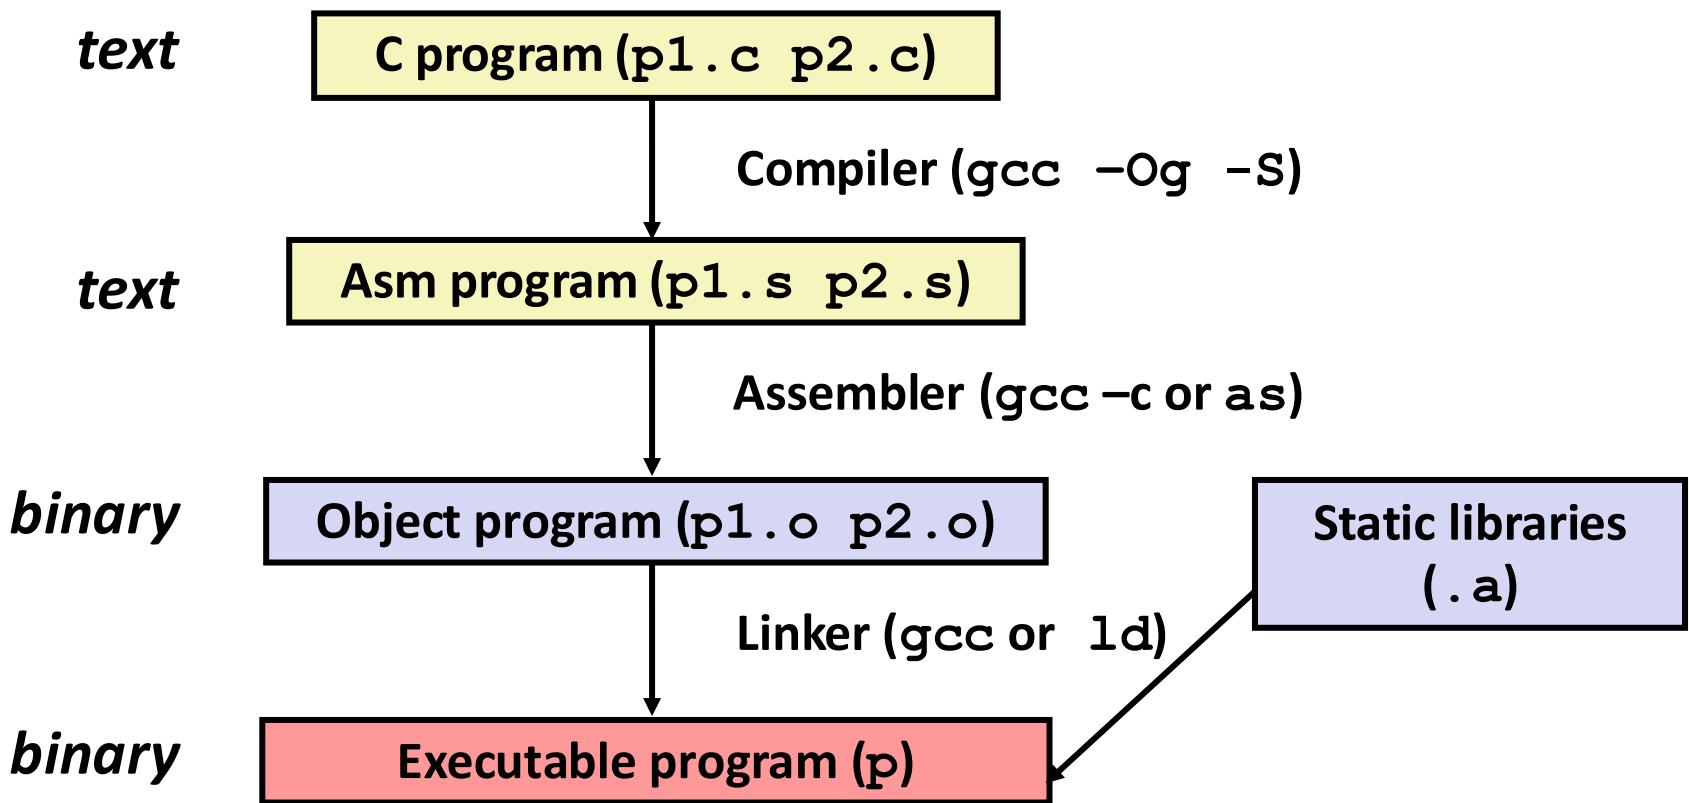
- **PC: Program counter**
  - Address of next instruction
  - Called “RIP” (x86-64)
- **Register file**
  - Heavily used program data
- **Condition codes**
  - Store status information about most recent arithmetic or logical operation
  - Used for conditional branching
- **Memory**
  - Byte addressable array
  - Code and user data
  - Stack to support procedures

# Where does it come from?

- **Assembly is usually generated by the compiler**
  - We can ask the compiler to show us the assembly
  - We can also generate assembly from machine code
- **The compiler runs a separate tool that generates machine code**
  - Machine code is just bytes in memory
- **Execution gives bytes their “types”**

# Turning C into Object Code

- Code in files `p1.c p2.c`
- Compile with command: `gcc -Og p1.c p2.c -o p`
  - Use debugging-friendly optimizations (`-Og`)
  - Put resulting binary in file `p`



# Compiling Into Assembly

## C Code (sum.c)

```
long plus(long x, long y);  
  
void sumstore(long x, long y,  
              long *dest)  
{  
    long t = plus(x, y);  
    *dest = t;  
}
```

## Generated x86-64 Assembly

```
sumstore:  
    pushq   %rbx  
    movq   %rdx, %rbx  
    call    plus  
    movq   %rax, (%rbx)  
    popq   %rbx  
    ret
```

Obtain (on shark machine) with command

```
gcc -Og -S sum.c
```

Produces file sum.s

**Warning:** Will get very different results on non-Shark machines (Andrew Linux, Mac OS-X, ...) due to different versions of gcc and different compiler settings.

# What it really looks like

```
.globl  sumstore
.type   sumstore, @function

sumstore:
.LFB35:
    .cfi_startproc
    pushq   %rbx
    .cfi_def_cfa_offset 16
    .cfi_offset 3, -16
    movq   %rdx, %rbx
    call   plus
    movq   %rax, (%rbx)
    popq   %rbx
    .cfi_def_cfa_offset 8
    ret
    .cfi_endproc

.LFE35:
    .size   sumstore, .-sumstore
```

# What it really looks like

```
.globl sumstore
.type sumstore, @function

sumstore:
.LFB35:
    .cfi_startproc
    pushq %rbx
    .cfi_def_cfa_offset 16
    .cfi_offset 3, -16
    movq %rdx, %rbx
    call plus
    movq %rax, (%rbx)
    popq %rbx
    .cfi_def_cfa_offset 8
    ret
    .cfi_endproc

.LFE35:
    .size sumstore, .-sumstore
```

Things that look weird  
and are preceded by a '**'**  
are generally directives.

```
sumstore:
    pushq %rbx
    movq %rdx, %rbx
    call plus
    movq %rax, (%rbx)
    popq %rbx
    ret
```

# Object Code

## Code for sumstore

0x0400595:

0x53

0x48

0x89

0xd3

0xe8

0xf2

0xff

0xff

0xff

0x48

0x89

0x03

0x5b

0xc3

- Total of 14 bytes
- Each instruction 1, 3, or 5 bytes
- Starts at address 0x0400595

### ■ Assembler

- Translates .s into .o
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing linkages between code in different files

### ■ Linker

- Resolves references between files
- Combines with static run-time libraries
  - E.g., code for **malloc**, **printf**
- Some libraries are *dynamically linked*
  - Linking occurs when program begins execution

# Machine Instruction Example

```
*dest = t;
```

## ■ C Code

- Store value **t** where designated by **dest**

```
movq %rax, (%rbx)
```

## ■ Assembly

- Move 8-byte value to memory
  - Quad words in x86-64 parlance
- Operands:
  - t:** Register **%rax**
  - dest:** Register **%rbx**
  - \*dest:** Memory **M[%rbx]**

```
0x40059e: 48 89 03
```

## ■ Object Code

- 3-byte instruction
- Stored at address **0x40059e**

# Disassembling Object Code

## Disassembled

```
0000000000400595 <sumstore>:  
 400595: 53                      push    %rbx  
 400596: 48 89 d3                mov     %rdx,%rbx  
 400599: e8 f2 ff ff ff         callq   400590 <plus>  
 40059e: 48 89 03                mov     %rax,(%rbx)  
 4005a1: 5b                      pop    %rbx  
 4005a2: c3                      retq
```

## ■ Disassembler

`objdump -d sum`

- Useful tool for examining object code
- Analyzes bit pattern of series of instructions
- Produces approximate rendition of assembly code
- Can be run on either a .out (complete executable) or .o file

# Alternate Disassembly

## Disassembled

```
Dump of assembler code for function sumstore:  
0x0000000000400595 <+0>: push    %rbx  
0x0000000000400596 <+1>: mov     %rdx,%rbx  
0x0000000000400599 <+4>: callq   0x400590 <plus>  
0x000000000040059e <+9>: mov     %rax,(%rbx)  
0x00000000004005a1 <+12>:pop    %rbx  
0x00000000004005a2 <+13>:retq
```

- Within gdb Debugger
    - Disassemble procedure
- gdb sum**
- disassemble sumstore**

# Alternate Disassembly

## Object Code

```
0x0400595:  
0x53  
0x48  
0x89  
0xd3  
0xe8  
0xf2  
0xff  
0xff  
0xff  
0x48  
0x89  
0x03  
0x5b  
0xc3
```

## Disassembled

```
Dump of assembler code for function sumstore:  
0x0000000000400595 <+0>: push    %rbx  
0x0000000000400596 <+1>: mov     %rdx,%rbx  
0x0000000000400599 <+4>: callq   0x400590 <plus>  
0x000000000040059e <+9>: mov     %rax,(%rbx)  
0x00000000004005a1 <+12>:pop    %rbx  
0x00000000004005a2 <+13>:retq
```

## ■ Within gdb Debugger

- Disassemble procedure

```
gdb sum
```

```
disassemble sumstore
```

- Examine the 14 bytes starting at sumstore

```
x/14xb sumstore
```

# What Can be Disassembled?

```
% objdump -d WINWORD.EXE

WINWORD.EXE:      file format pei-i386

No symbols in "WINWORD.EXE".
Disassembly of section .text:

30001000 <.text>:
30001000:
30001001:
30001003:
30001005:
3000100a:
```

**Reverse engineering forbidden by  
Microsoft End User License Agreement**

- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

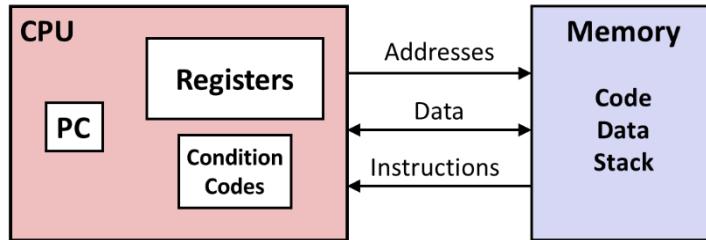
# Appendix

# Levels of Abstraction

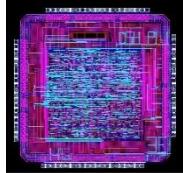
C programmer

```
#include <stdio.h>
int main(){
    int i, n = 10, t1 = 0, t2 = 1, nxt;
    for (i = 1; i <= n; ++i){
        printf("%d, ", t1);
        nxt = t1 + t2;
        t1 = t2;
        t2 = nxt; }
    return 0; }
```

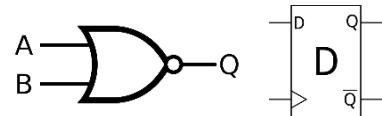
Assembly programmer



Computer Designer



Gates, clocks, circuit layout, ...

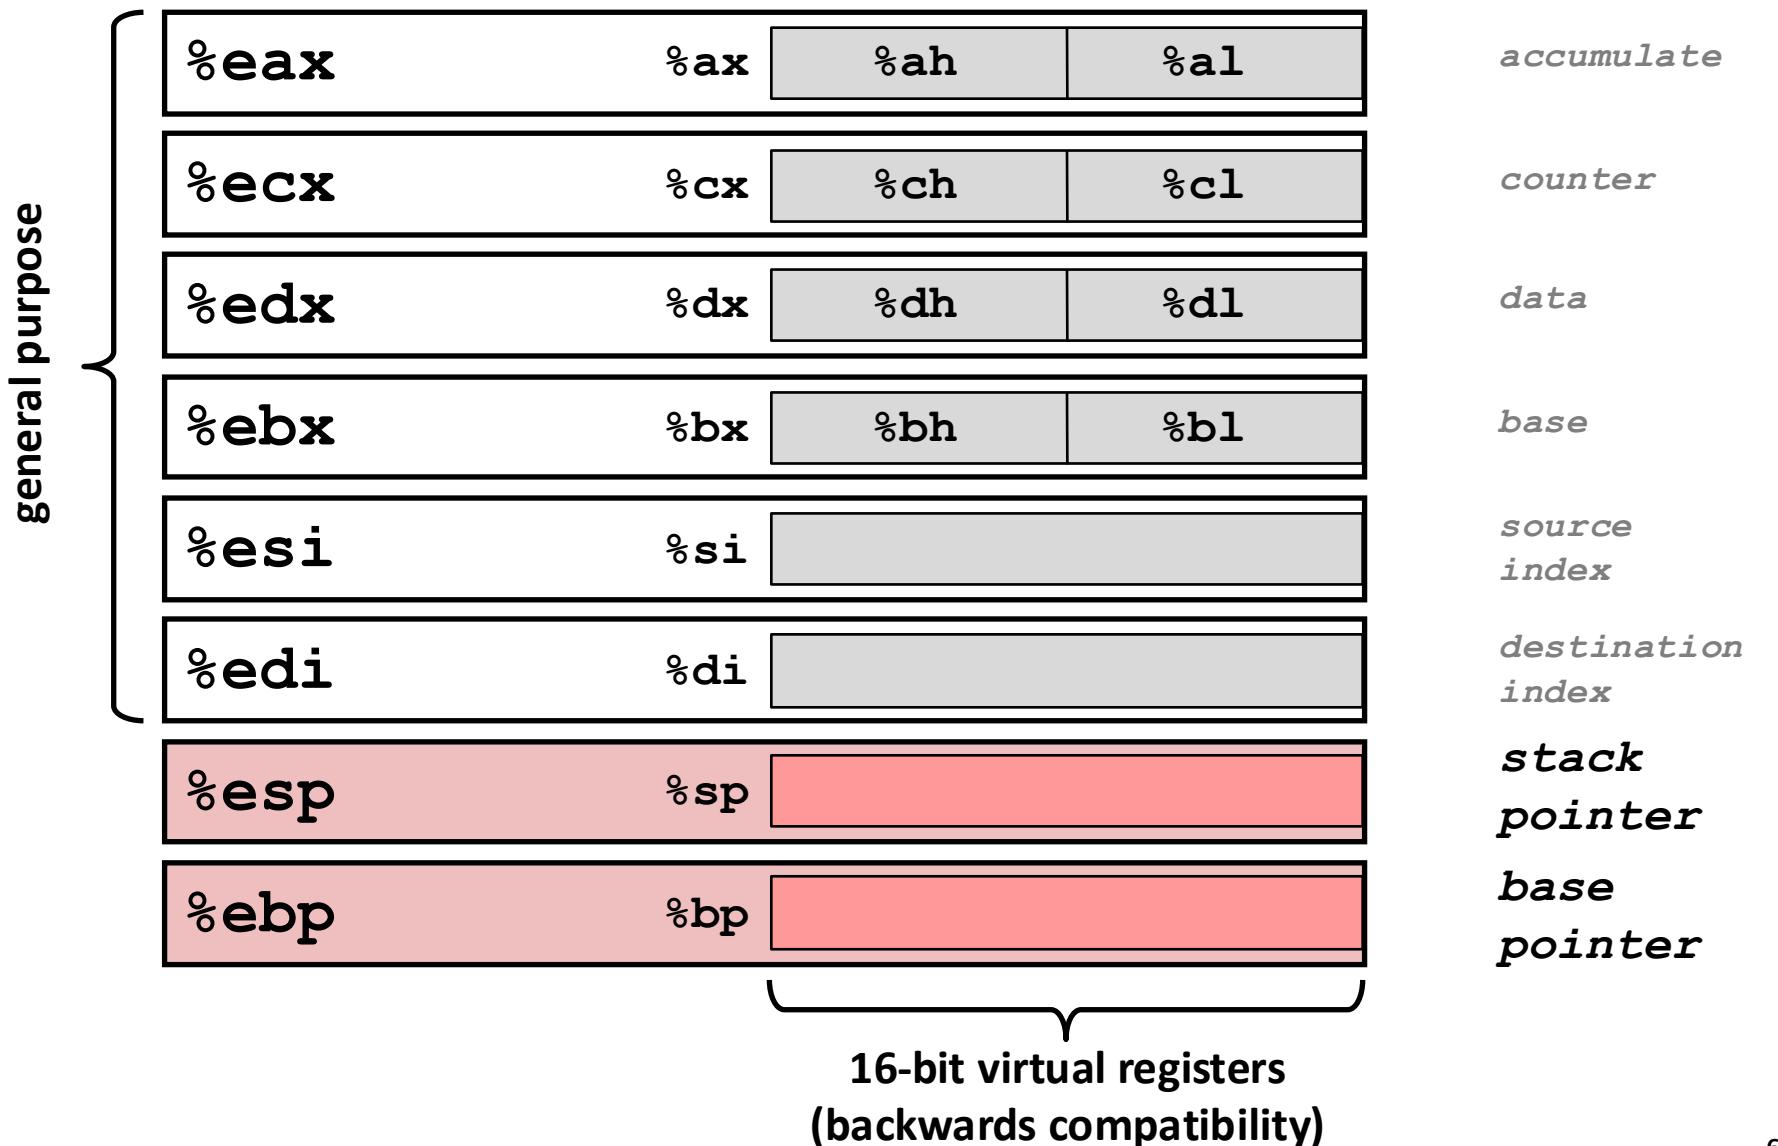


# x86-64 Integer Registers

%rax	%eax	%r8	%r8d
%rbx	%ebx	%r9	%r9d
%rcx	%ecx	%r10	%r10d
%rdx	%edx	%r11	%r11d
%rsi	%esi	%r12	%r12d
%rdi	%edi	%r13	%r13d
%rsp	%esp	%r14	%r14d
%rbp	%ebp	%r15	%r15d

- Can reference low-order 4 bytes (also low-order 1 & 2 bytes)
- Not part of memory (or cache)

# Some History: IA32 Registers

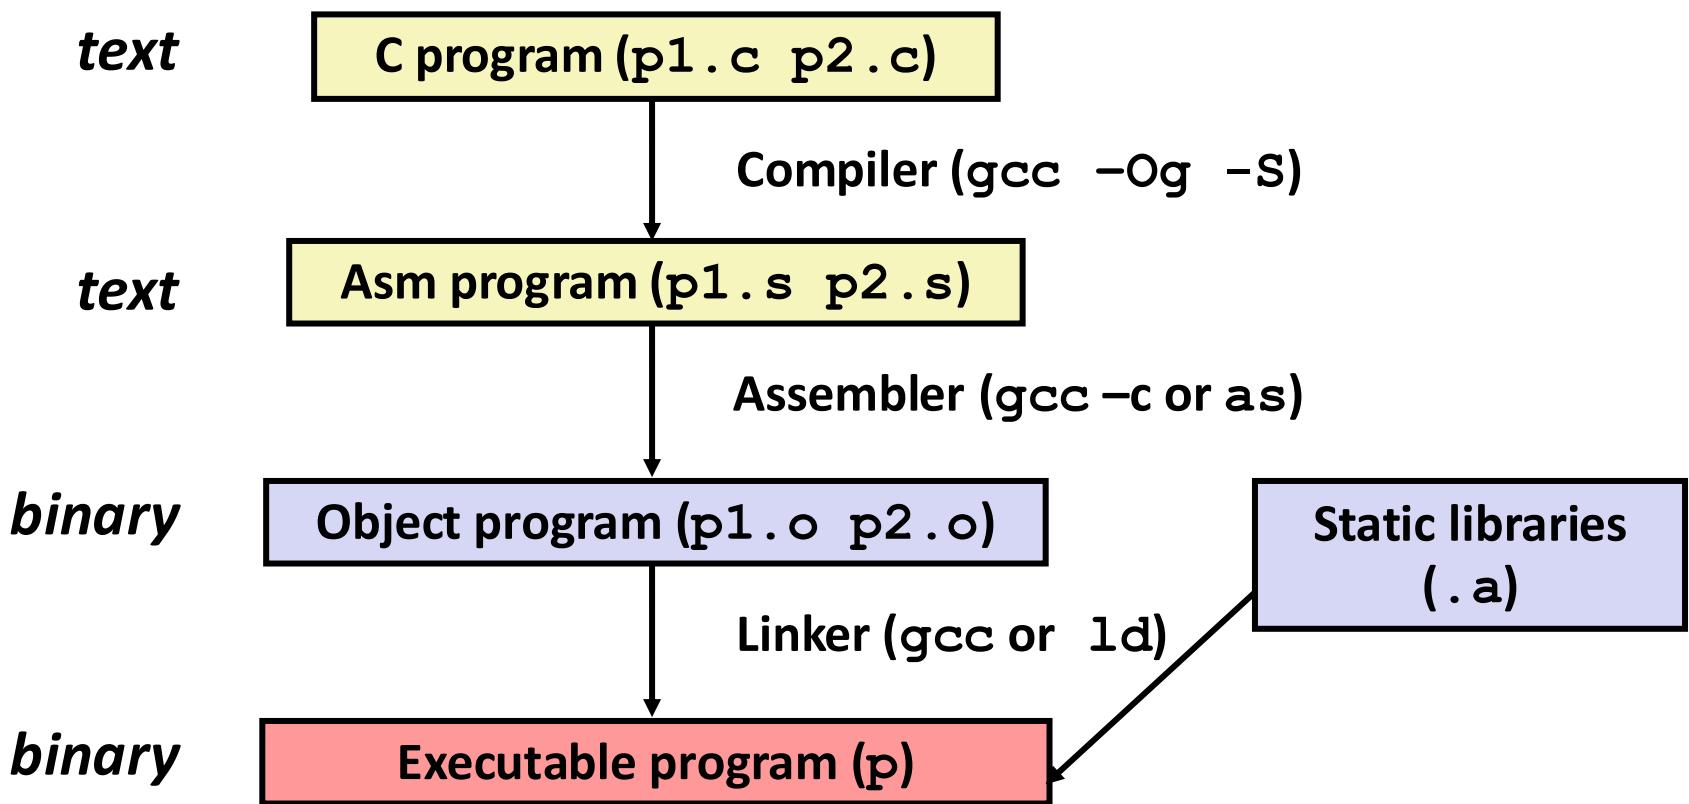


# Today: Machine Programming I: Basics

- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- C, assembly, machine code

# Turning C into Object Code

- Code in files `p1.c p2.c`
- Compile with command: `gcc -Og p1.c p2.c -o p`
  - Use debugging-friendly optimizations (`-Og`)
  - Put resulting binary in file `p`



# Compiling Into Assembly

## C Code (sum.c)

```
long plus(long x, long y);  
  
void sumstore(long x, long y,  
              long *dest)  
{  
    long t = plus(x, y);  
    *dest = t;  
}
```

## Generated x86-64 Assembly

```
sumstore:  
    pushq   %rbx  
    movq   %rdx, %rbx  
    call    plus  
    movq   %rax, (%rbx)  
    popq   %rbx  
    ret
```

Obtain (on shark machine) with command

```
gcc -Og -S sum.c
```

Produces file sum.s

**Warning:** Will get very different results on non-Shark machines (Andrew Linux, Mac OS-X, ...) due to different versions of gcc and different compiler settings.

# What it really looks like

```
.globl  sumstore
.type   sumstore, @function

sumstore:
.LFB35:
    .cfi_startproc
    pushq   %rbx
    .cfi_def_cfa_offset 16
    .cfi_offset 3, -16
    movq   %rdx, %rbx
    call   plus
    movq   %rax, (%rbx)
    popq   %rbx
    .cfi_def_cfa_offset 8
    ret
    .cfi_endproc

.LFE35:
    .size   sumstore, .-sumstore
```

# What it really looks like

```
.globl sumstore
.type sumstore, @function

sumstore:
.LFB35:
    .cfi_startproc
    pushq %rbx
    .cfi_def_cfa_offset 16
    .cfi_offset 3, -16
    movq %rdx, %rbx
    call plus
    movq %rax, (%rbx)
    popq %rbx
    .cfi_def_cfa_offset 8
    ret
    .cfi_endproc

.LFE35:
    .size sumstore, .-sumstore
```

Things that look weird  
and are preceded by a '**'**  
are generally directives.

```
sumstore:
    pushq %rbx
    movq %rdx, %rbx
    call plus
    movq %rax, (%rbx)
    popq %rbx
    ret
```

# Object Code

## Code for sumstore

0x0400595:

0x53

0x48

0x89

0xd3

0xe8

0xf2

0xff

0xff

0xff

0x48

0x89

0x03

0x5b

0xc3

- Total of 14 bytes
- Each instruction 1, 3, or 5 bytes
- Starts at address 0x0400595

### ■ Assembler

- Translates .s into .o
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing linkages between code in different files

### ■ Linker

- Resolves references between files
- Combines with static run-time libraries
  - E.g., code for **malloc**, **printf**
- Some libraries are *dynamically linked*
  - Linking occurs when program begins execution

# Machine Instruction Example

```
*dest = t;
```

## ■ C Code

- Store value **t** where designated by **dest**

```
movq %rax, (%rbx)
```

## ■ Assembly

- Move 8-byte value to memory
  - Quad words in x86-64 parlance
- Operands:
  - t:** Register **%rax**
  - dest:** Register **%rbx**
  - \*dest:** Memory **M[%rbx]**

```
0x40059e: 48 89 03
```

## ■ Object Code

- 3-byte instruction
- Stored at address **0x40059e**

# Disassembling Object Code

## Disassembled

```
0000000000400595 <sumstore>:  
 400595: 53                      push    %rbx  
 400596: 48 89 d3                mov     %rdx,%rbx  
 400599: e8 f2 ff ff ff         callq   400590 <plus>  
 40059e: 48 89 03                mov     %rax,(%rbx)  
 4005a1: 5b                      pop    %rbx  
 4005a2: c3                      retq
```

## ■ Disassembler

`objdump -d sum`

- Useful tool for examining object code
- Analyzes bit pattern of series of instructions
- Produces approximate rendition of assembly code
- Can be run on either a .out (complete executable) or .o file

# Alternate Disassembly

## Disassembled

```
Dump of assembler code for function sumstore:  
0x0000000000400595 <+0>: push    %rbx  
0x0000000000400596 <+1>: mov     %rdx,%rbx  
0x0000000000400599 <+4>: callq   0x400590 <plus>  
0x000000000040059e <+9>: mov     %rax,(%rbx)  
0x00000000004005a1 <+12>:pop    %rbx  
0x00000000004005a2 <+13>:retq
```

- Within gdb Debugger
    - Disassemble procedure
- gdb sum**
- disassemble sumstore**

# Alternate Disassembly

## Object Code

```
0x0400595:  
0x53  
0x48  
0x89  
0xd3  
0xe8  
0xf2  
0xff  
0xff  
0xff  
0x48  
0x89  
0x03  
0x5b  
0xc3
```

## Disassembled

```
Dump of assembler code for function sumstore:  
0x0000000000400595 <+0>: push    %rbx  
0x0000000000400596 <+1>: mov     %rdx,%rbx  
0x0000000000400599 <+4>: callq   0x400590 <plus>  
0x000000000040059e <+9>: mov     %rax,(%rbx)  
0x00000000004005a1 <+12>:pop    %rbx  
0x00000000004005a2 <+13>:retq
```

## ■ Within gdb Debugger

- Disassemble procedure

```
gdb sum
```

```
disassemble sumstore
```

- Examine the 14 bytes starting at sumstore

```
x/14xb sumstore
```

# What Can be Disassembled?

```
% objdump -d WINWORD.EXE

WINWORD.EXE:      file format pei-i386

No symbols in "WINWORD.EXE".
Disassembly of section .text:

30001000 <.text>:
30001000:
30001001:
30001003:
30001005:
3000100a:
```

**Reverse engineering forbidden by  
Microsoft End User License Agreement**

- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

# Machine Programming I: Summary

- **History of Intel processors and architectures**
  - Evolutionary design leads to many quirks and artifacts
- **C, assembly, machine code**
  - New forms of visible state: program counter, registers, ...
  - Compiler must transform statements, expressions, procedures into low-level instruction sequences
- **Assembly Basics: Registers, operands, move**
  - The x86-64 move instructions cover wide range of data movement forms
- **Arithmetic**
  - C compiler will figure out different instruction combinations to carry out computation

# History: Machine Programming I: Basics

- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- C, assembly, machine code

# Intel x86 Processors

- Dominate laptop/desktop/server market
- Evolutionary design
  - Backwards compatible up until 8086, introduced in 1978
  - Added more features as time goes on
    - Now 3 volumes, about 5,000 pages of documentation
- Complex instruction set computer (**CISC**)
  - Many different instructions with many different formats
    - But, only small subset encountered with Linux programs
  - Hard to match performance of Reduced Instruction Set Computers (RISC)
  - But, Intel has done just that!
    - In terms of speed. Less so for low power.

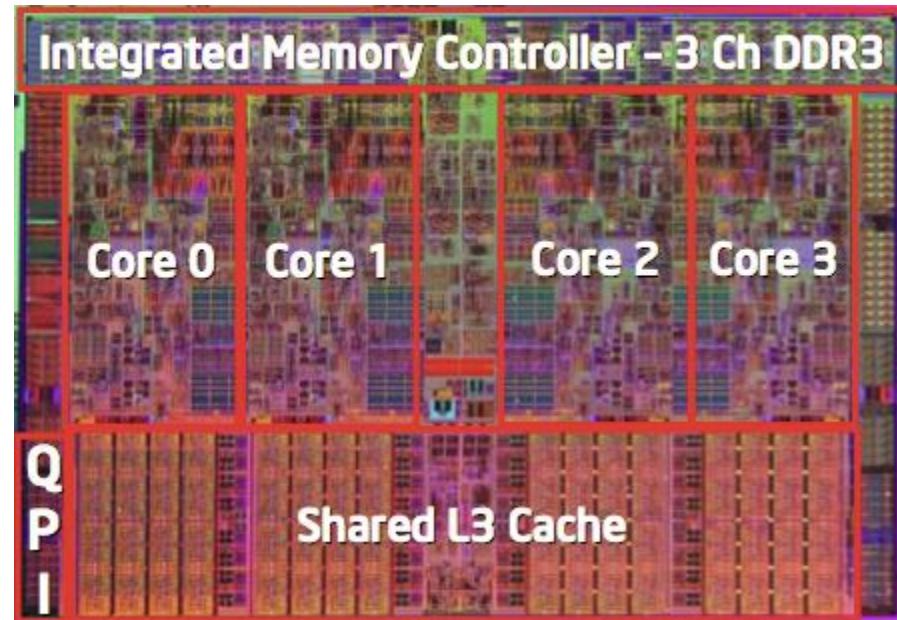
# Intel x86 Evolution: Milestones

<i>Name</i>	<i>Date</i>	<i>Transistors</i>	<i>MHz</i>
■ <b>8086</b>	<b>1978</b>	<b>29K</b>	<b>5-10</b>
■	▪ First 16-bit Intel processor. Basis for IBM PC & DOS		
	▪ 1MB address space		
■ <b>386</b>	<b>1985</b>	<b>275K</b>	<b>16-33</b>
■	▪ First 32 bit Intel processor , referred to as IA32		
	▪ Added “flat addressing”, capable of running Unix		
■ <b>Pentium 4E</b>	<b>2004</b>	<b>125M</b>	<b>2800-3800</b>
■	▪ First 64-bit Intel x86 processor, referred to as x86-64		
■ <b>Core 2</b>	<b>2006</b>	<b>291M</b>	<b>1060-3333</b>
■	▪ First multi-core Intel processor		
■ <b>Core i7</b>	<b>2008</b>	<b>731M</b>	<b>1600-4400</b>
■	▪ Four cores (our shark machines)		

# Intel x86 Processors, cont.

## ■ Machine Evolution

■ 386	1985	0.3M
■ Pentium	1993	3.1M
■ Pentium/MMX	1997	4.5M
■ PentiumPro	1995	6.5M
■ Pentium III	1999	8.2M
■ Pentium 4	2000	42M
■ Core 2 Duo	2006	291M
■ Core i7	2008	731M
■ Core i7 Skylake	2015	1.9B



## ■ Added Features

- Instructions to support multimedia operations
- Instructions to enable more efficient conditional operations
- Transition from 32 bits to 64 bits
- More cores

# Intel x86 Processors, cont.

## ■ Past Generations

		Process technology
■	1 <sup>st</sup> Pentium Pro	1995
■	1 <sup>st</sup> Pentium III	1999
■	1 <sup>st</sup> Pentium 4	2000
■	1 <sup>st</sup> Core 2 Duo	2006

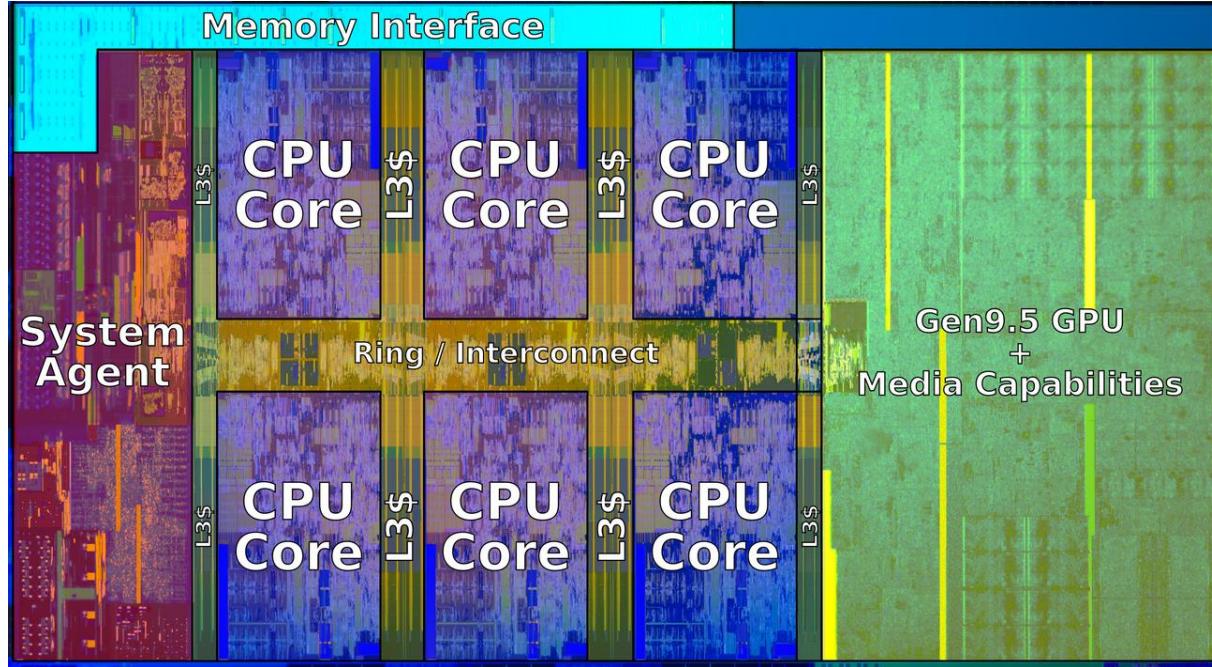
## ■ Recent & Upcoming Generations

1.	Nehalem	2008	45 nm
2.	Sandy Bridge	2011	32 nm
3.	Ivy Bridge	2012	22 nm
4.	Haswell	2013	22 nm
5.	Broadwell	2014	14 nm
6.	Skylake	2015	14 nm
7.	Kaby Lake	2016	14 nm
8.	Coffee Lake	2017	14 nm
9.	Cannon Lake	2018	10 nm
10.	Ice Lake	2019	10 nm
11.	Tiger Lake	2020	10 nm
12.	Alder Lake	2022	“intel 7” (10nm++++)

Process technology dimension  
= width of narrowest wires  
(10 nm ≈ 100 atoms wide)

(But this is changing now.)

# 2018 State of the Art: Coffee Lake



## ■ Mobile Model: Core i7

- 2.2-3.2 GHz
- 45 W

## ■ Desktop Model: Core i7

- Integrated graphics
- 2.4-4.0 GHz
- 35-95 W

## ■ Server Model: Xeon E

- Integrated graphics
- Multi-socket enabled
- 3.3-3.8 GHz
- 80-95 W

# x86 Clones: Advanced Micro Devices (AMD)

## ■ Historically

- AMD has followed just behind Intel
- A little bit slower, a lot cheaper

## ■ Then

- Recruited top circuit designers from Digital Equipment Corp. and other downward trending companies
- Built Opteron: tough competitor to Pentium 4
- Developed x86-64, their own extension to 64 bits

## ■ Recent Years

- Intel got its act together
  - 1995-2011: Lead semiconductor “fab” in world
  - 2018: #2 largest by \$\$ (#1 is Samsung)
  - 2019: reclaimed #1
- AMD fell behind: Spun off GlobalFoundries
- 2019-20: Pulled ahead! Used TSMC for part of fab
- 2022: Intel re-took the lead

# Intel's 64-Bit History

- **2001: Intel Attempts Radical Shift from IA32 to IA64**
  - Totally different architecture (Itanium)
  - Executes IA32 code only as legacy
  - Performance disappointing
- **2003: AMD Steps in with Evolutionary Solution**
  - x86-64 (now called “AMD64”)
- **Intel Felt Obligated to Focus on IA64**
  - Hard to admit mistake or that AMD is better
- **2004: Intel Announces EM64T extension to IA32**
  - Extended Memory 64-bit Technology
  - Almost identical to x86-64!
- **All but low-end x86 processors support x86-64**
  - But, lots of code still runs in 32-bit mode

# Our Coverage

## ■ IA32

- The traditional x86
- For 15/18-213: RIP, Summer 2015

## ■ x86-64

- The standard
- shark> gcc hello.c
- shark> gcc -m64 hello.c

## ■ Presentation

- Book covers x86-64
- Web aside on IA32
- We will only cover x86-64